



Operating Systems

Dr. Shu Yin

Part II: Process Management

- Processes
- Threads
- Process Synchronization
- CPU Scheduling
- Deadlocks



Goals

- Processes
- Threads



Thread

- A fundamental unit of CPU utilization that forms the basis of multithreaded computer systems



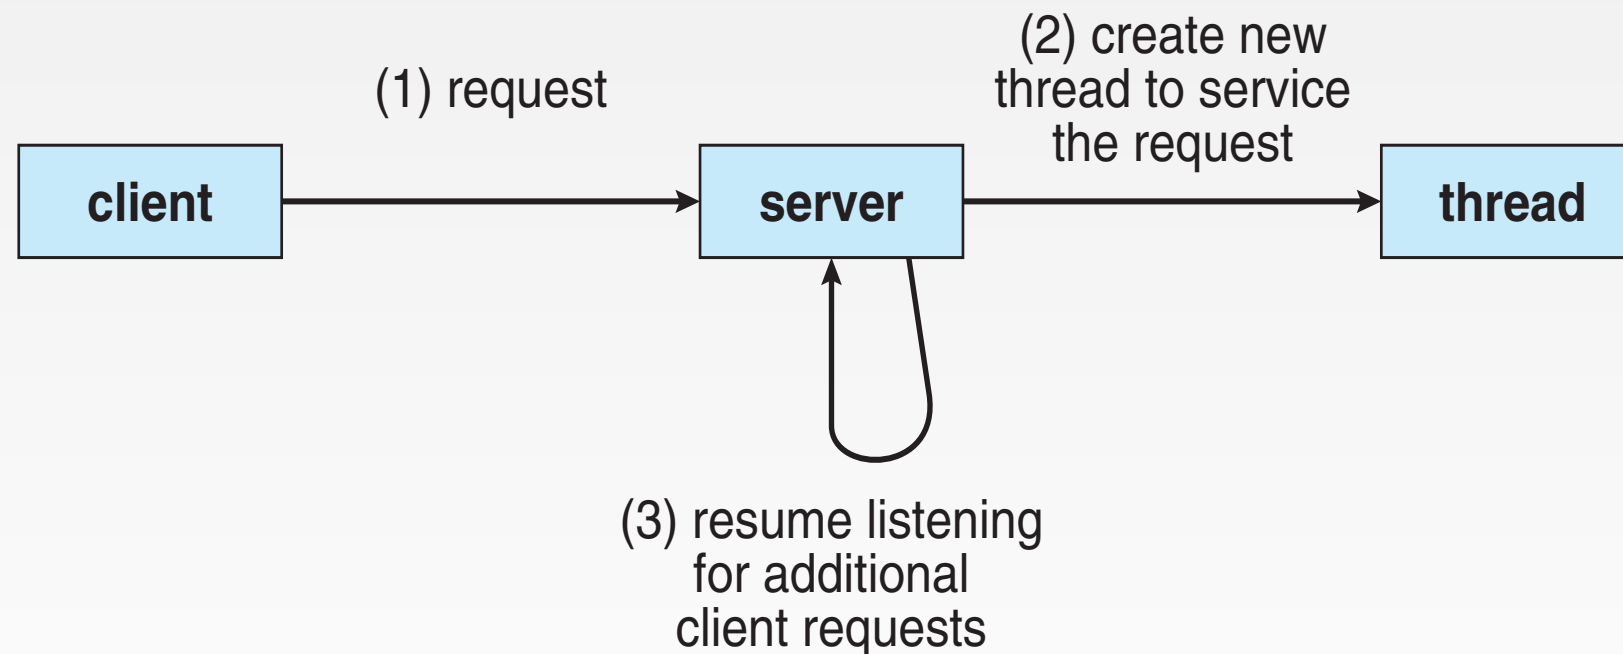
Motivation

- Most modern apps are multithreaded
- Threads run within application
- Multiple tasks with the app can be implemented by separate threads
 - update display
 - fetch data
 - spell checking
 - answer a network request

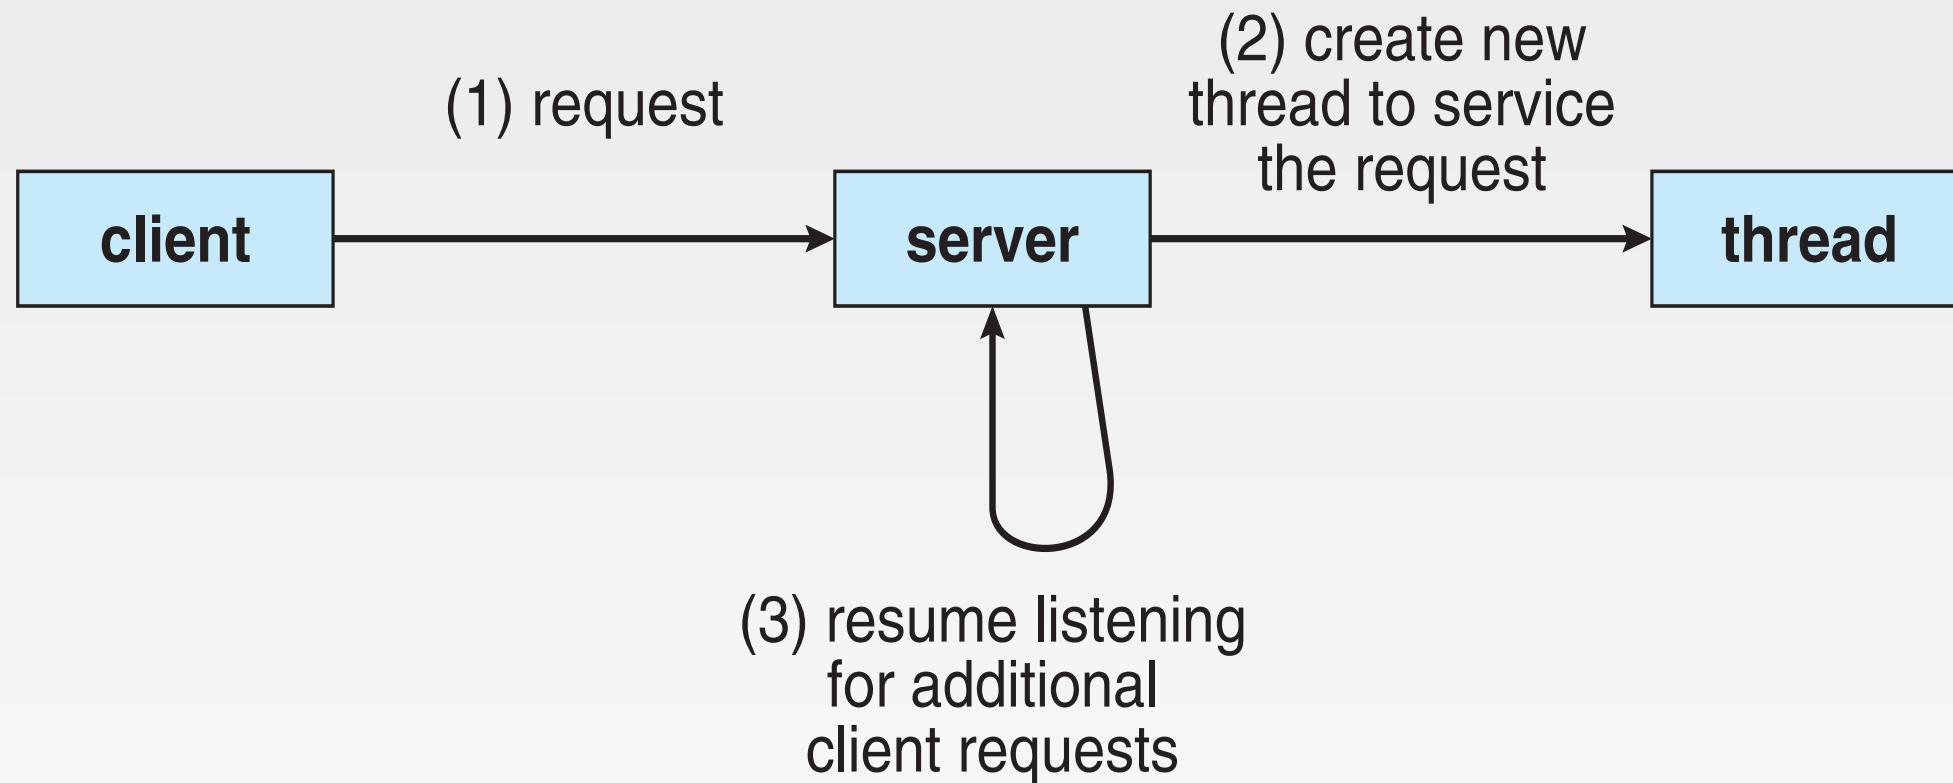


Motivation (cont.)

- Can simplify code, increase efficiency
- Kernels are generally multithreaded



Multithreaded Server Architecture



Benefit

- Responsiveness
- Resource sharing
- Economy
- Scalability



Multicore programming

- Multicore systems putting pressure on programmers, challenges include:
 - Dividing activities
 - Balance
 - Data splitting
 - Data dependency
 - Testing and debugging



Multicore programming (cont.)

- Parallelism
 - A system can perform more than one task simultaneously
- Concurrency
 - Supports more than one task making progress
 - in single core, scheduler provide concurrency




Multicore programming (cont.)

- Types of parallelism
 - Data parallelism
 - subset of same data across multiple cores
 - same operation on each
 - Task parallelism
 - threads across cores, each thread performing unique operation





Recall: Concurrency vs. Parallelism



time

concurrent execution on single-core system



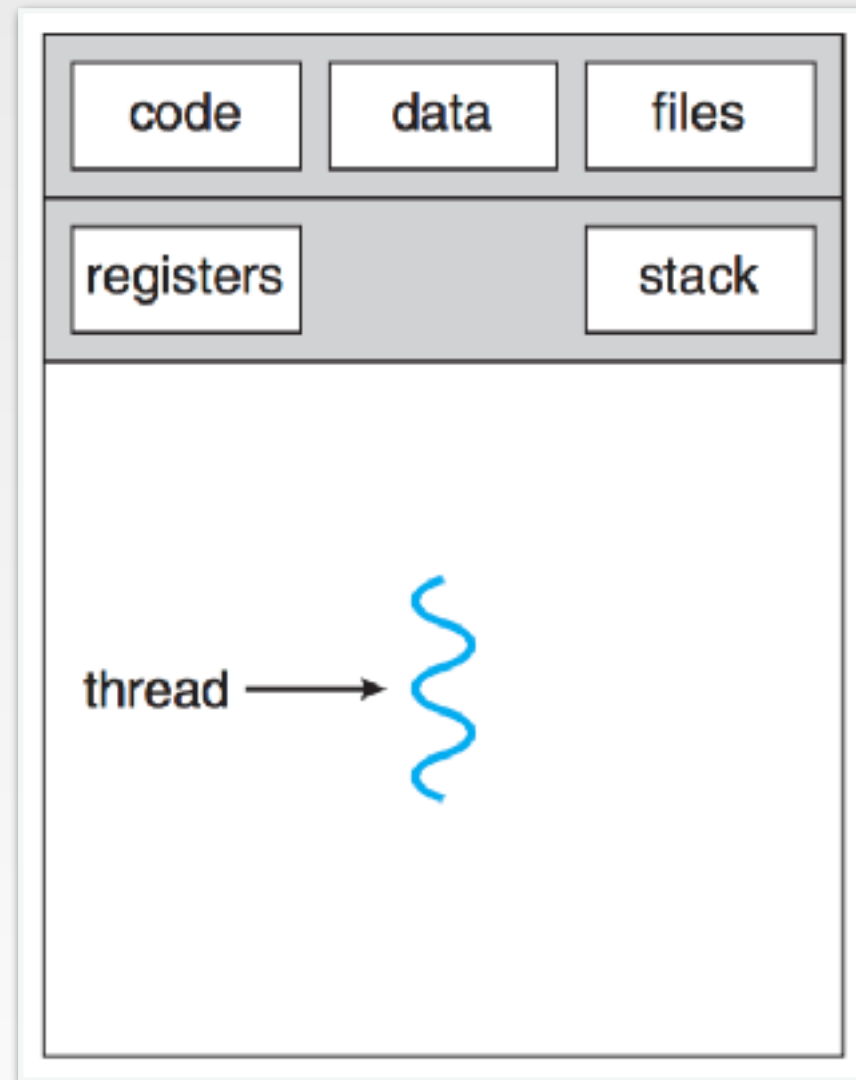
time

parallelism on a multi-core system



上海科技大学
ShanghaiTech University

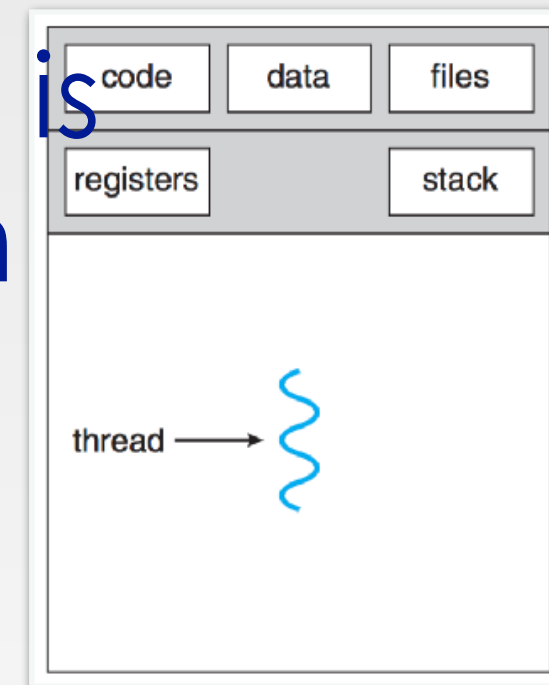
Recall: Traditional UNIX Process



Single-threaded Process

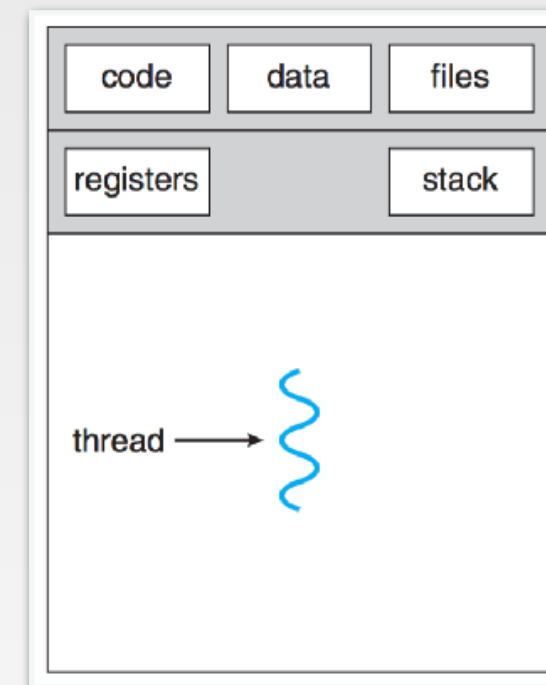
Single Threaded Process (cont.)

- Process: OS abstraction of what is needed to run a single program
 - “Heavyweight Process”
 - No concurrency
- Two Parts:
 - Active Part: Sequential program execution stream
 - Passive Part: Protected resources



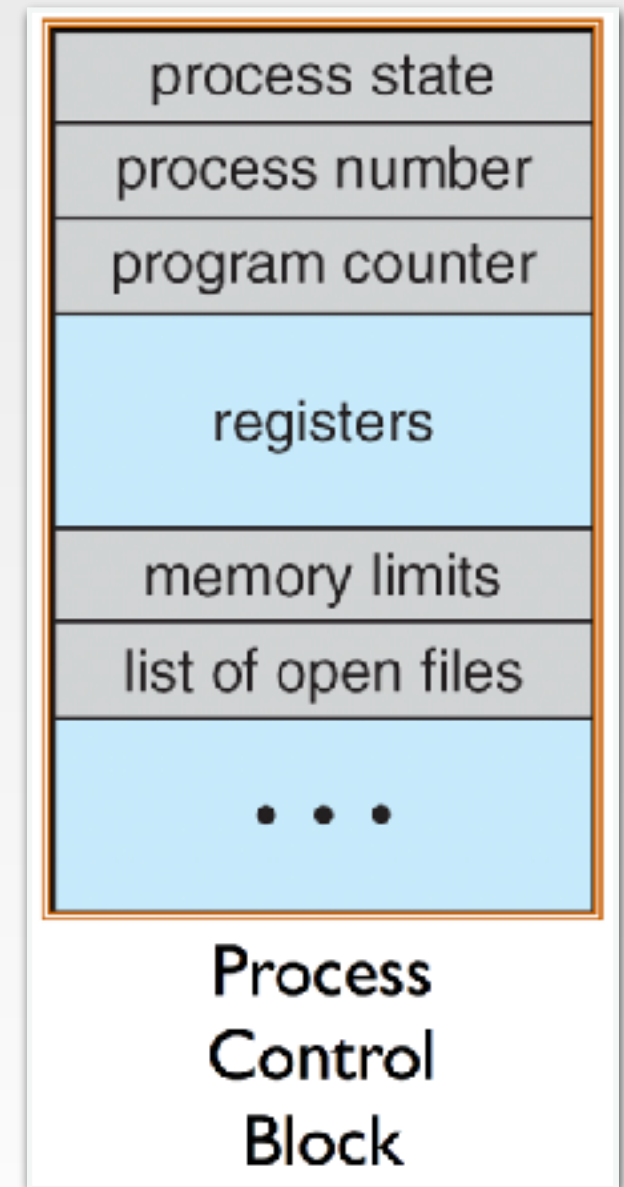
Single Threaded Process (cont.)

- Active Part
 - Code executed as a stream of execution (i.e. thread)
 - Includes state of CPU registers
- Passive Part
 - Main memory state (contents of Address Space)
 - I/O state (i.e. file descriptors)



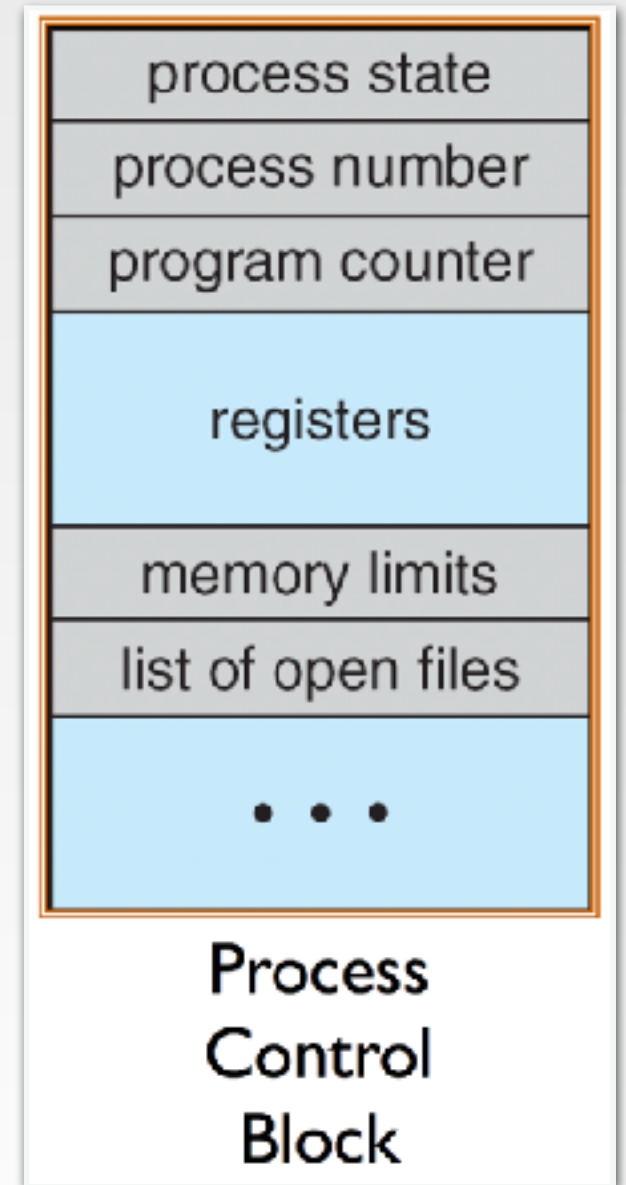
How do we Multiplex Processes?

- The current state of process held in a PCB
 - A “snapshot” of the execution and protection environment
 - Only one PCB active at a time



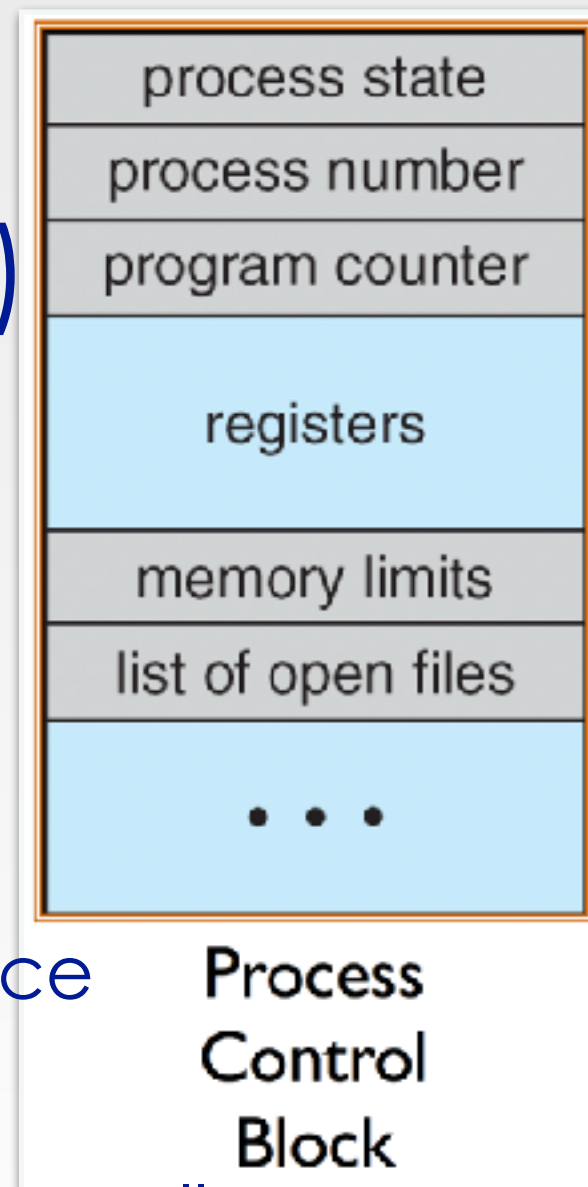
Multiplex Processes(cont.)

- Give out CPU time to different processes (Scheduling)
 - Only one process “running” at a time
 - Give more time to important processes

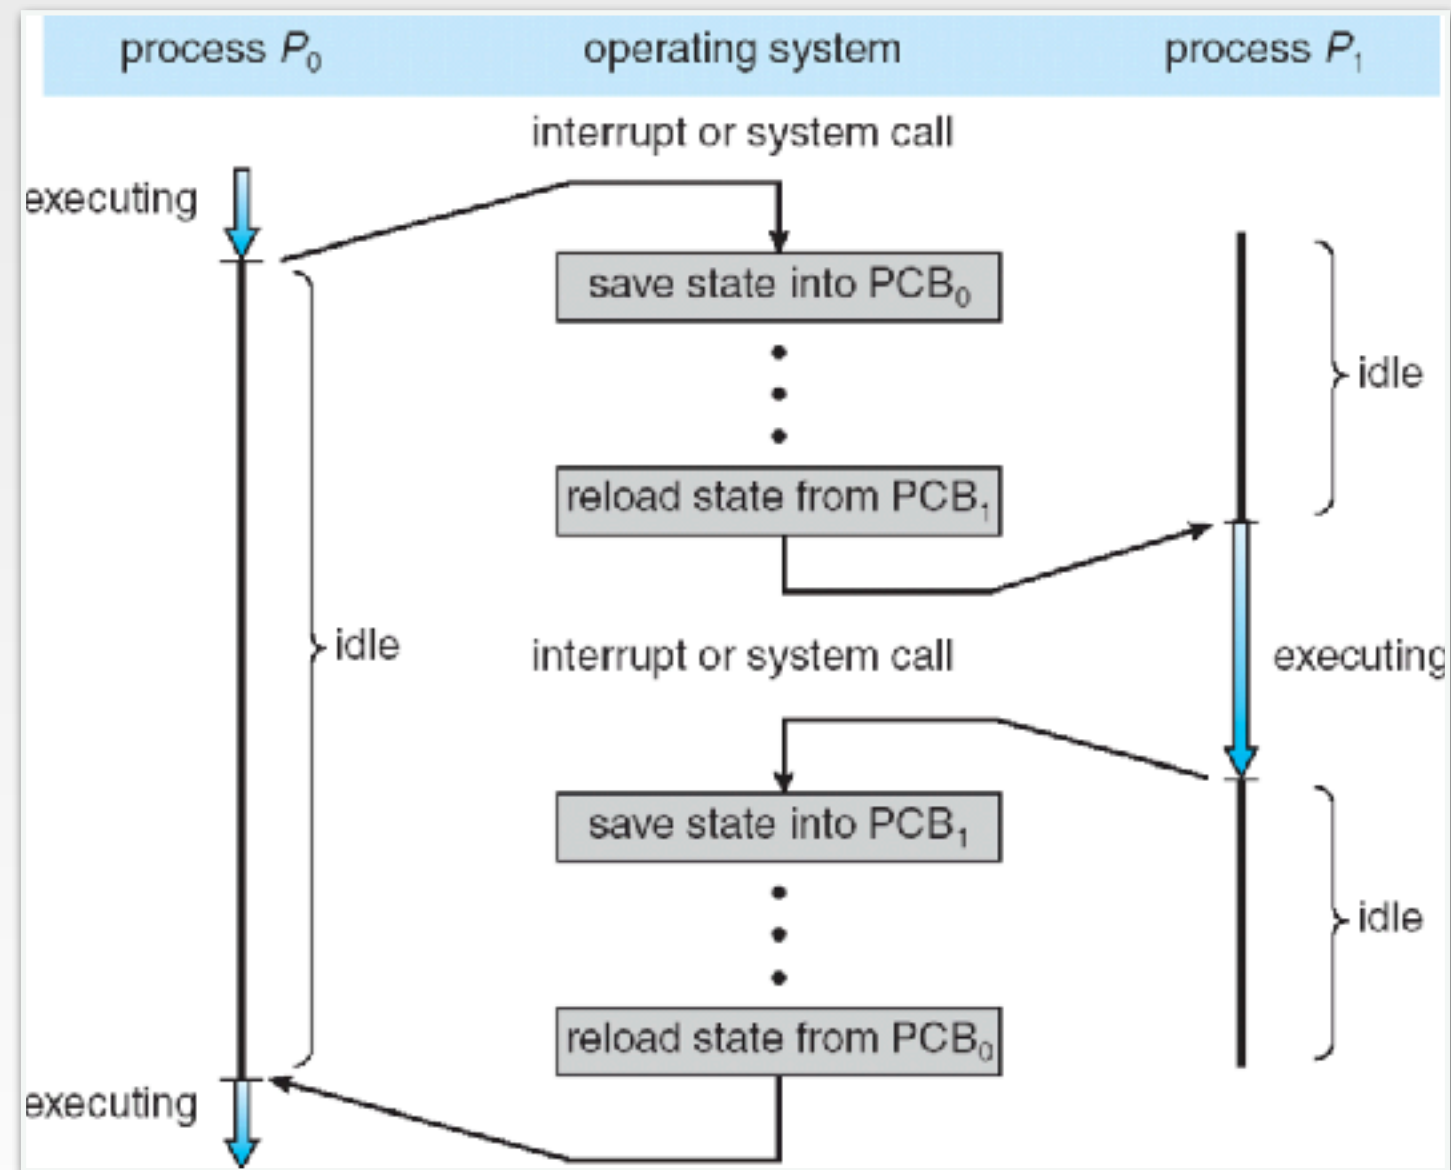


Multiplex Processes(cont.)

- Give pieces of resources to different processes (Protection)
 - Controlled access to non-CPU resources
 - Example mechanisms:
 - Memory Mapping
 - Give each process their own address space
 - Kernel/User duality
 - Arbitrary multiplexing of I/O through system

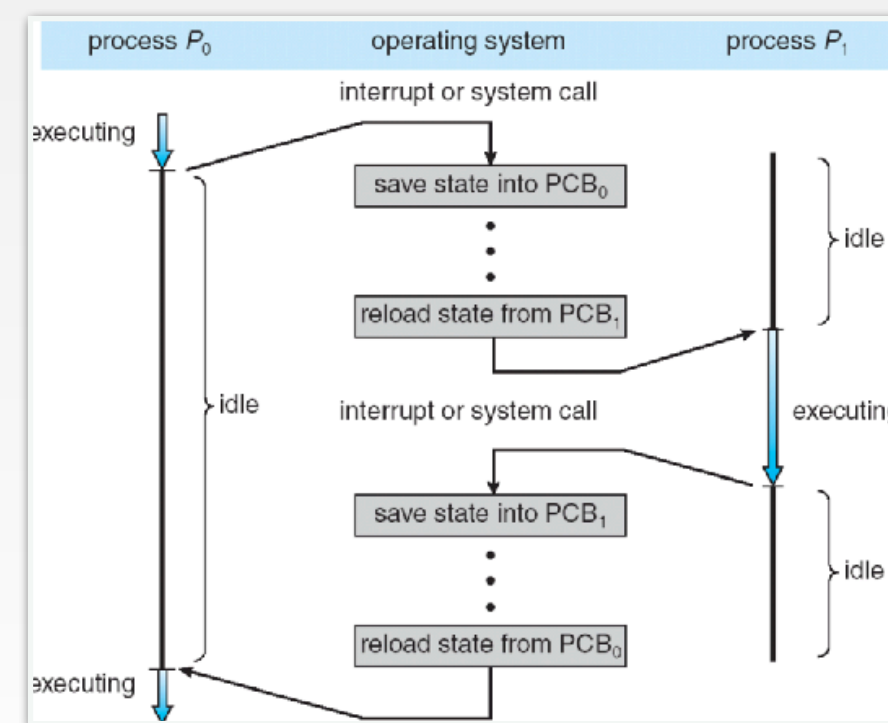


CPU Switch from Process A to B

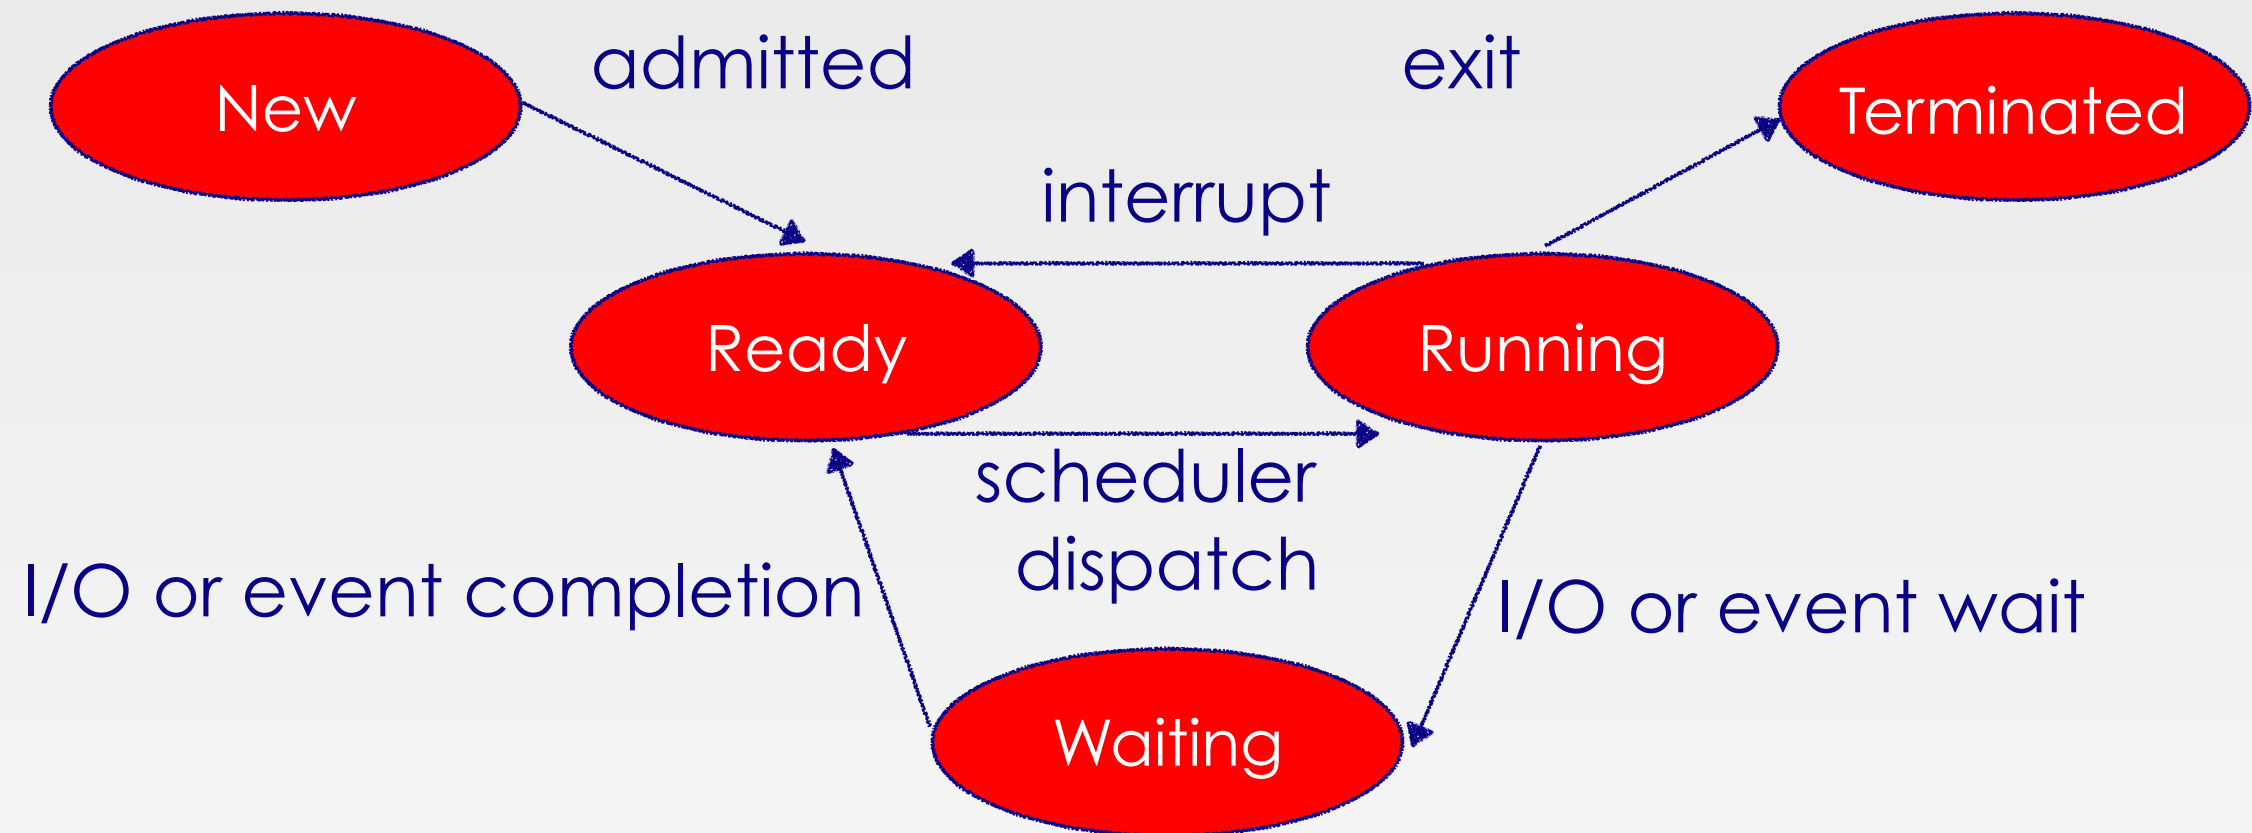


Switch from Process A to B (cont.)

- Context Switch
- Code executed in kernel is overhead
 - Overhead sets minimum practical switching time
 - Less overhead with SMT/HT, but contention for resources

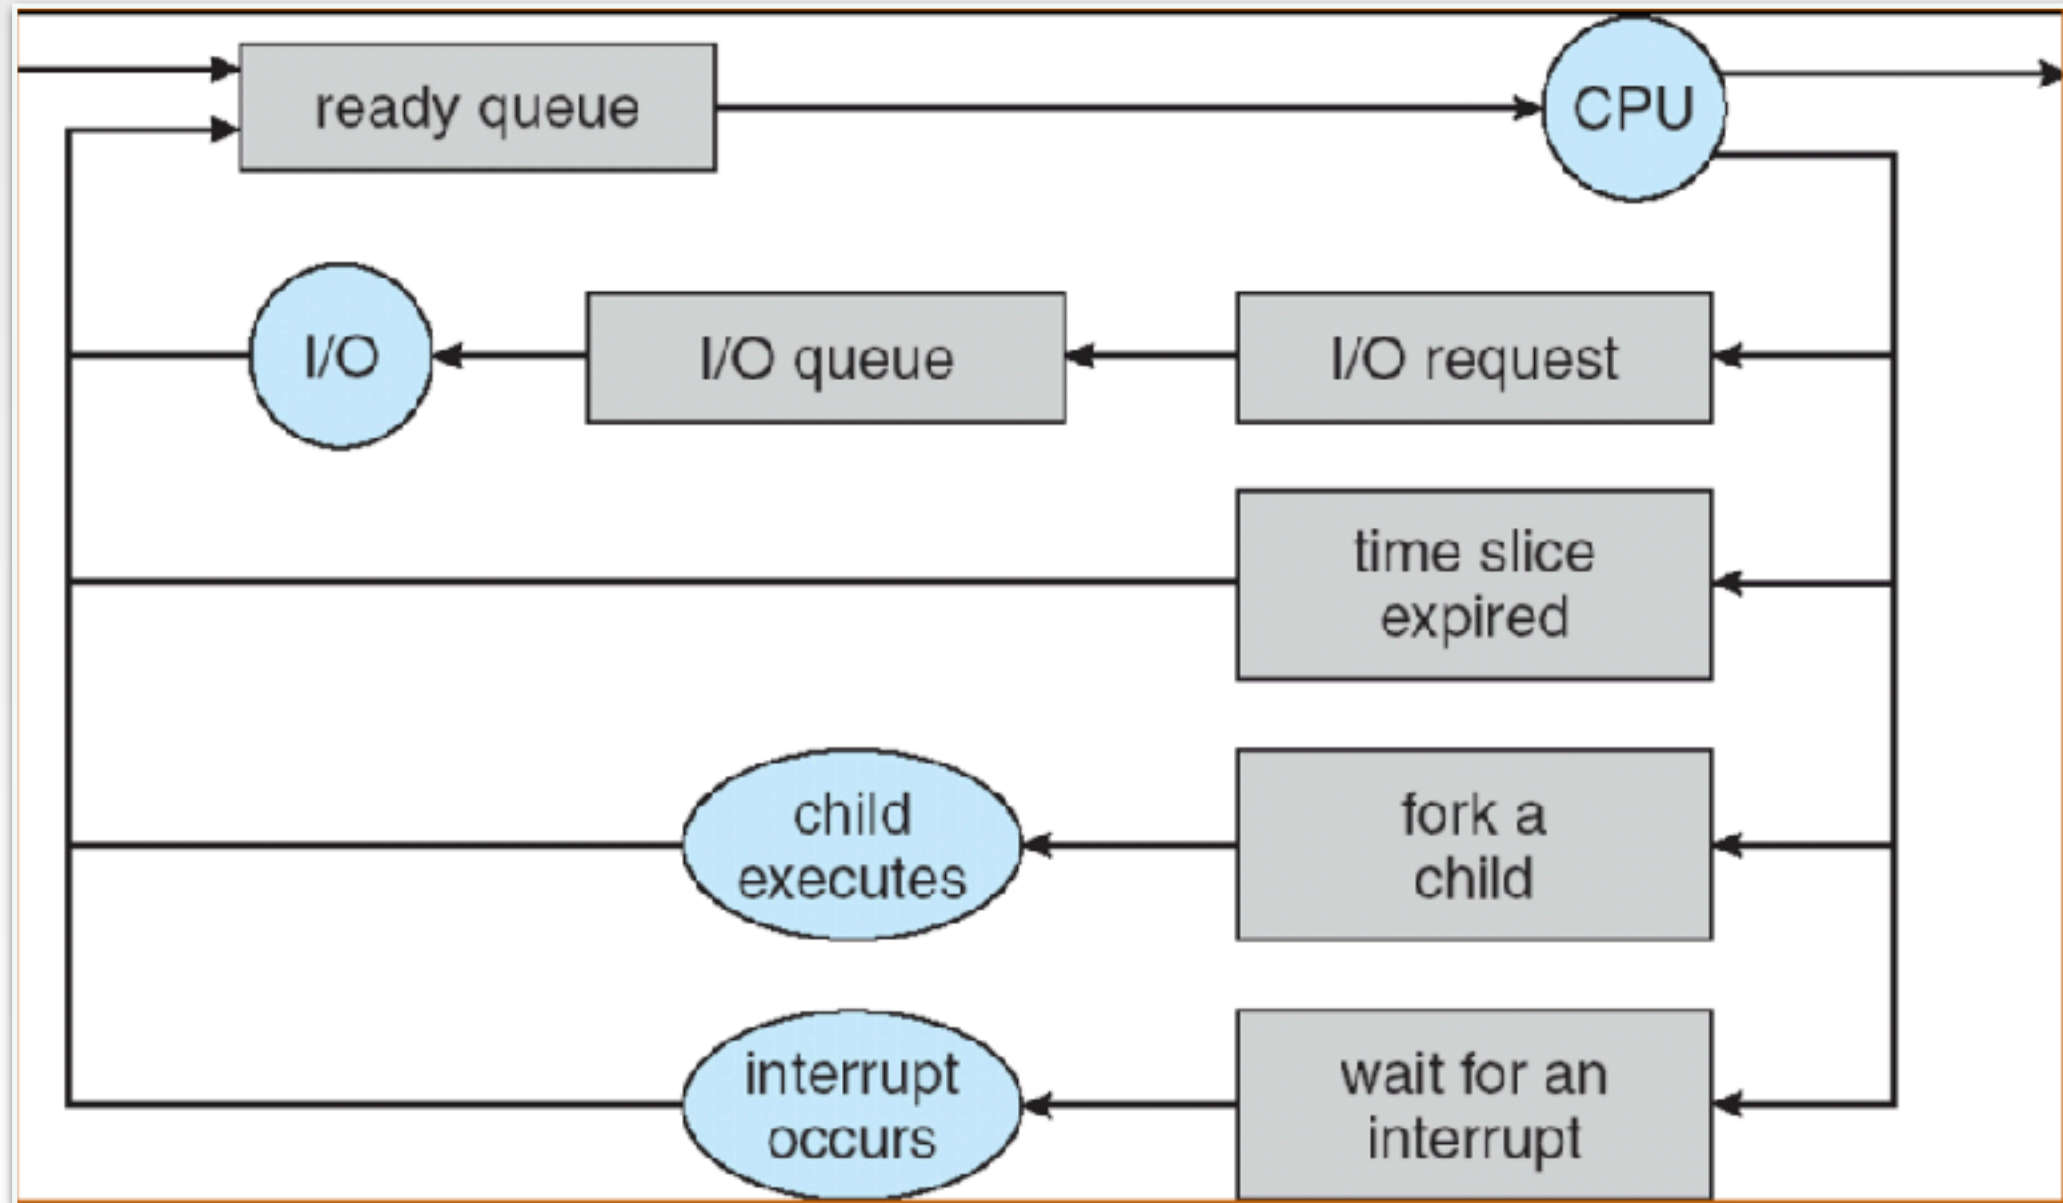


Lifecycle of a Process



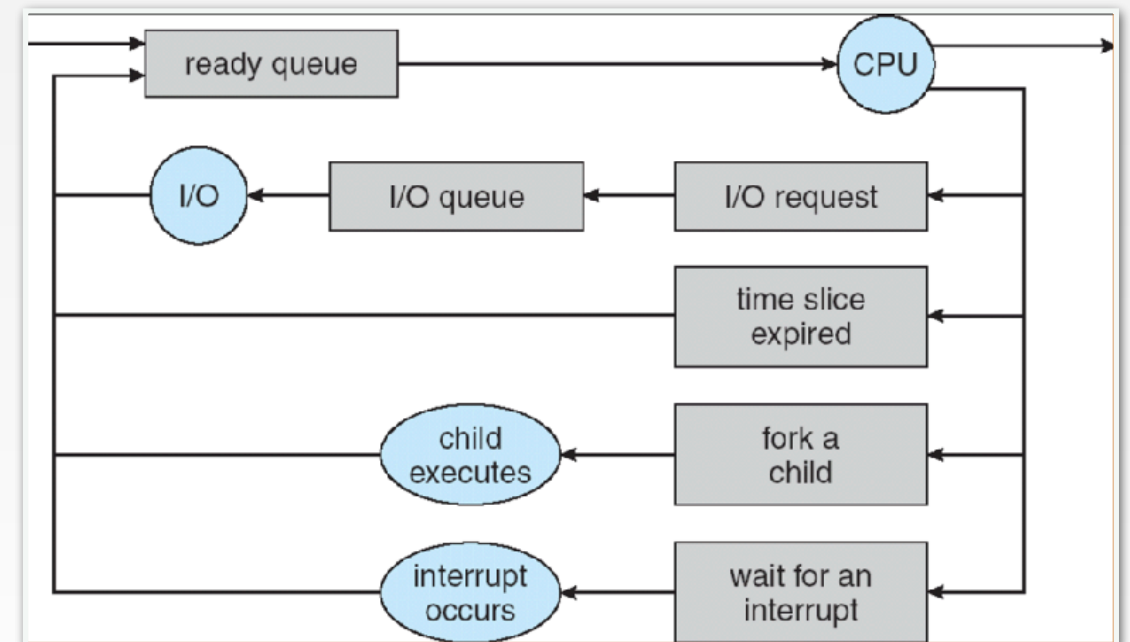
- As a process executes, it changes state
 - New, Ready, Running, Waiting, Terminated

Process Scheduling

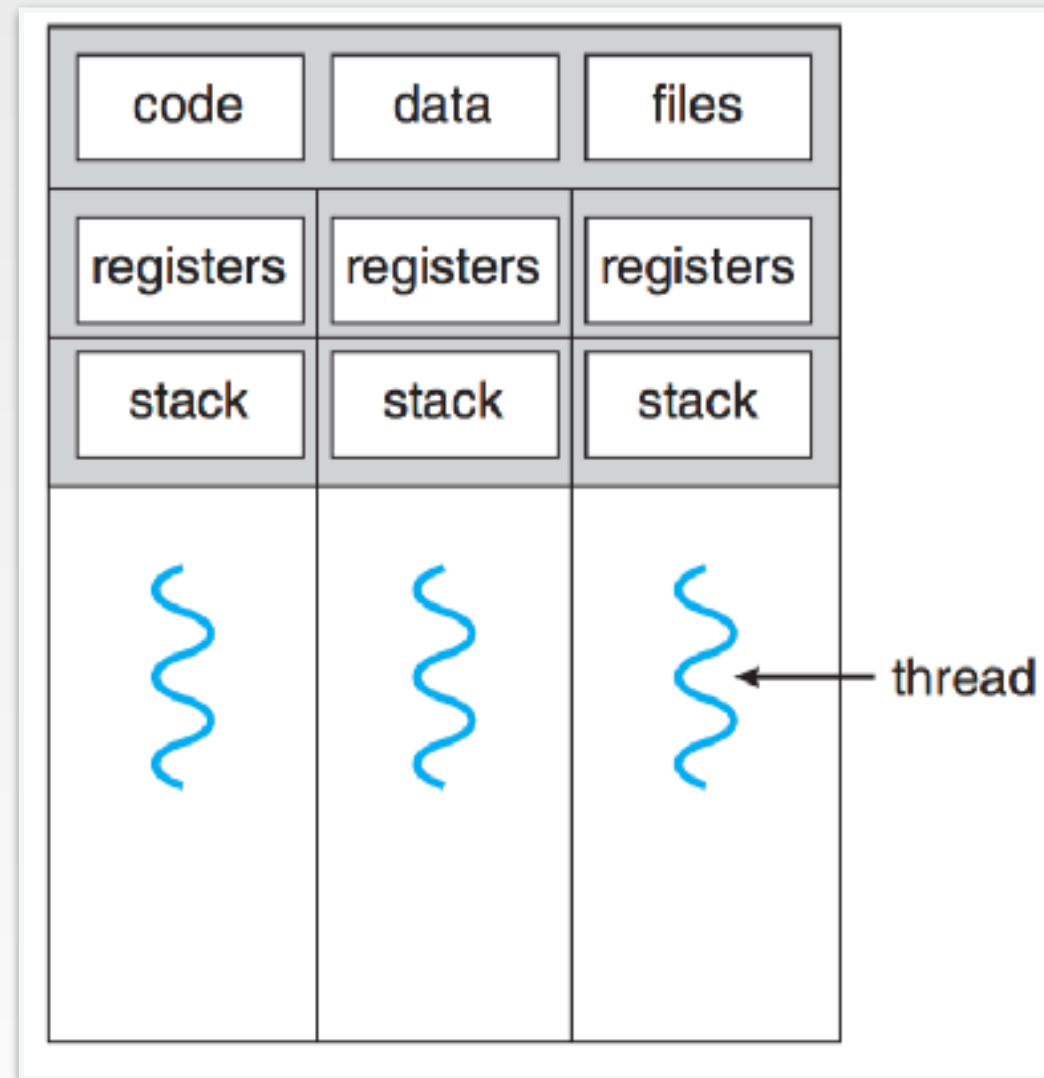


Process Scheduling (cont.)

- PCBs move from queue to queue as they change state
 - Decisions about which order to remove from queues are Scheduling decisions
 - Many Algorithms



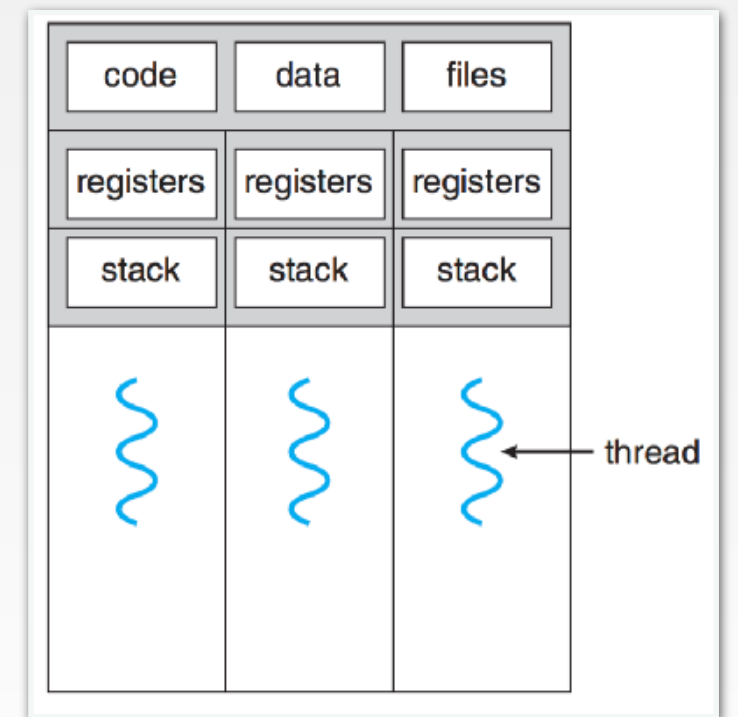
Modern Process with Threads



Multithreaded Process

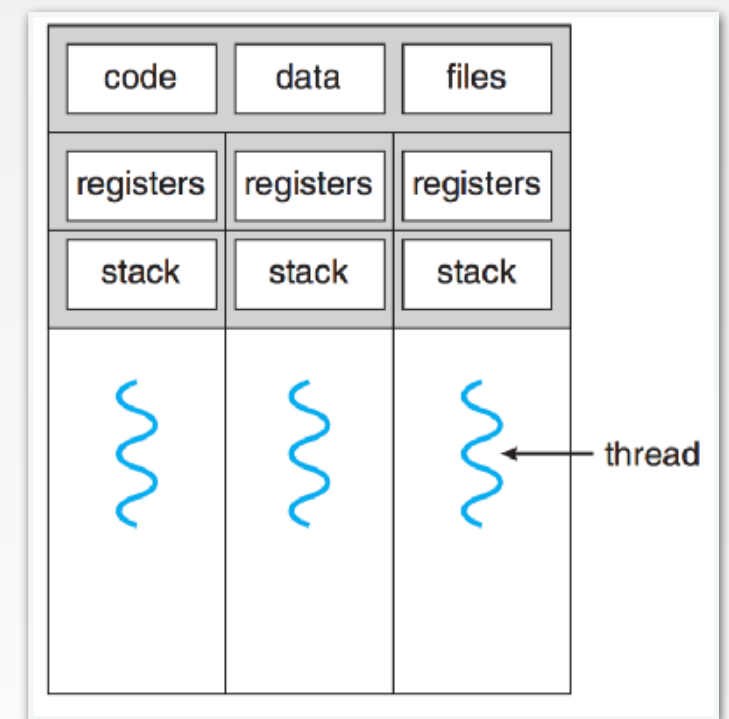
Modern Process with Threads

- Thread: a sequential execution stream within process (“Lightweight Process”)
 - Process still contains a single Address Space
 - No protection between threads



Modern Process with Threads

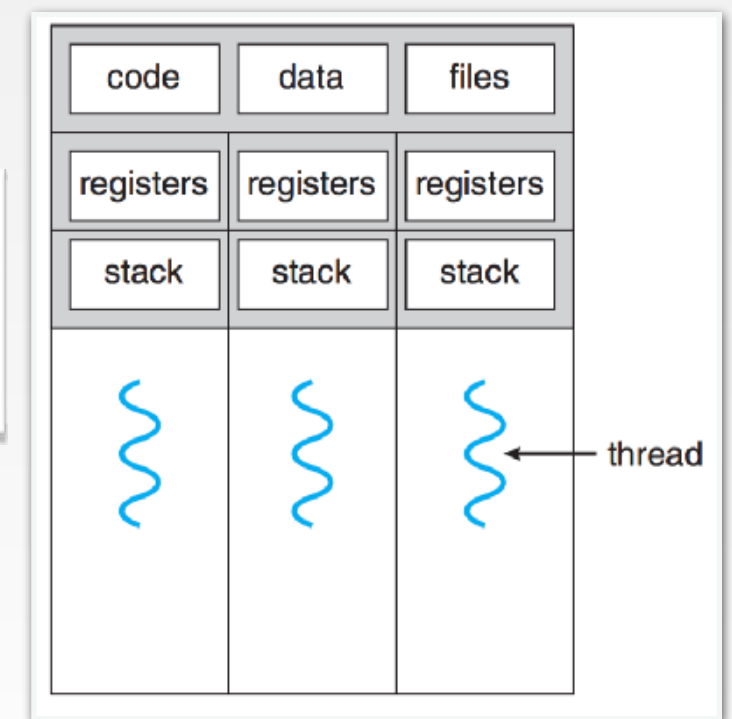
- Multithreading: a single program, a number of different concurrent activities
 - Sometimes called multitasking



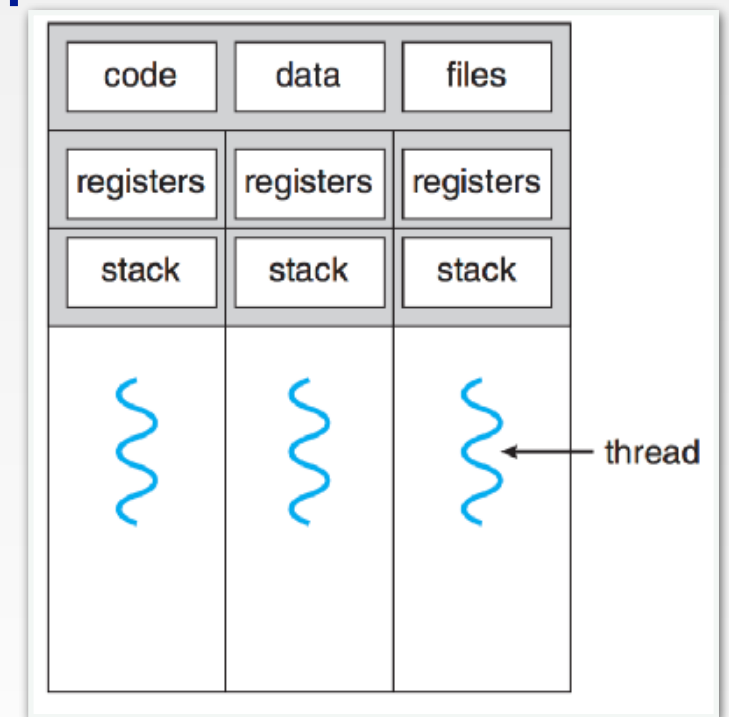
Modern Process with Threads

- Multithreading: a single program, a number of different concurrent activities
 - Sometimes called multitasking

Why separate the concept of a thread from that of a process?




- Discuss the “thread” part of a process
 - Concurrency
- Separate from the “address space”
 - Protection
- Heavyweight process is
 - Process with one thread

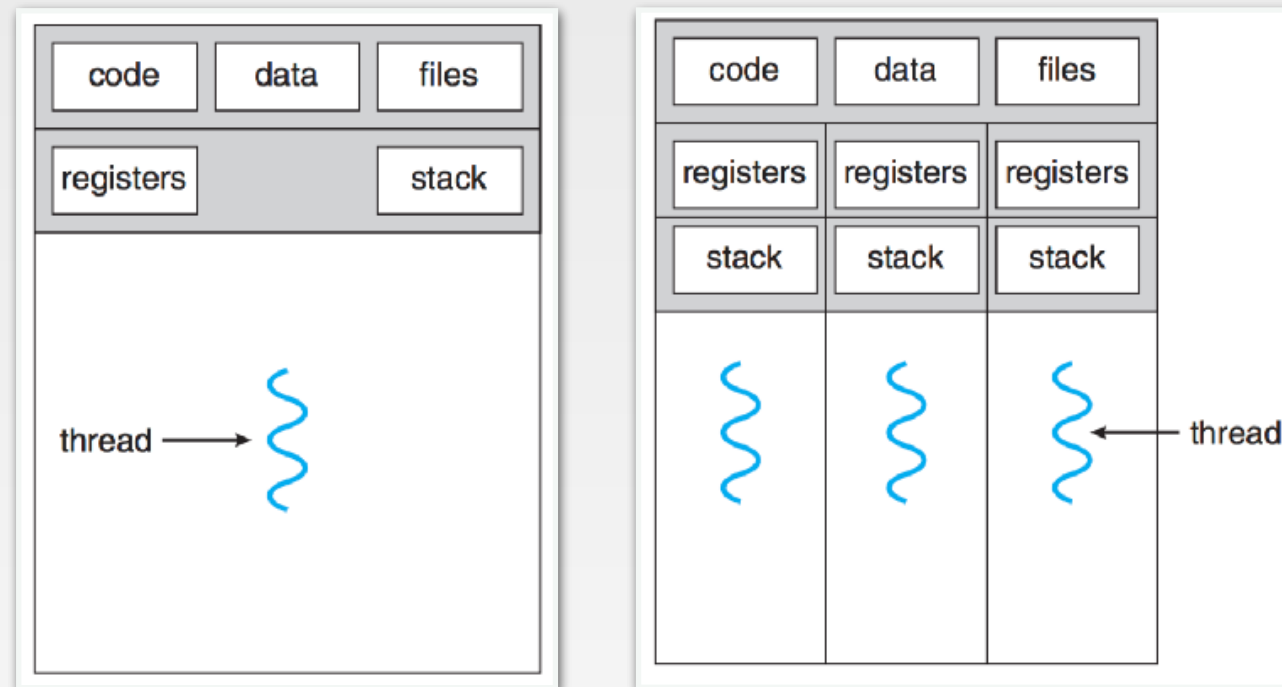


Why separate the concept of a thread from that of a process?

- Discuss the “thread” part of a process
 - Concurrency
- Separate from the “address space”
 - Protection
- Heavyweight process is
 - Process with one thread

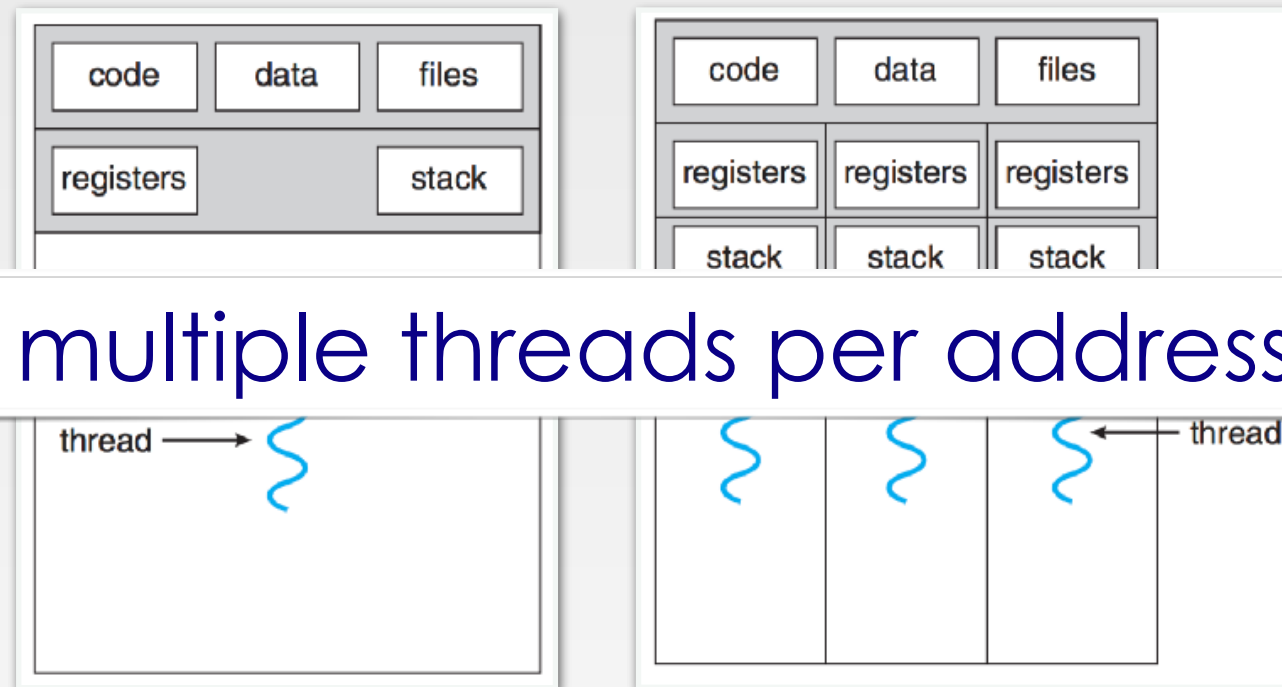


Single vs. Multithreaded Processes



- Thread encapsulate concurrency (Active part)
- Address space encapsulate protection (Passive part)
 - Keeps buggy program from trashing the system

Single vs. Multithreaded Processes



Why have multiple threads per address space?

- Thread encapsulate concurrency (Active part)
- Address space encapsulate protection (Passive part)
 - Keeps buggy program from trashing the system

In general, Threads

- In a multiple threaded task, while one server thread is blocked and waiting, a second thread in the same task can run
 - Cooperation of multiple threads in the same job confers higher throughput and improved performance
 - Apps that require sharing a common buffer



Threads (cont.)

- Threads provide a mechanism that allows sequential processes to make blocking system calls
- While also achieving parallelism
- Threads share CPU, only one thread can run at a time



Thread State

- State shared by all threads in process/
address space
 - Content of memory (global variables, heap)
 - I/O state (file descriptors, network connections)
- State “private” to each thread
 - Kept in Thread Control Block (TCB)
 - CPU registers (including PC)
 - Execution Stack



Thread State (cont.)

- Execution Stack
 - Parameters, temporary variables
 - Return PCs are kept while called procedures are executing
- Four states:
 - ready, blocked, running, terminated



Shared vs. Per-Thread State

Shared State

Heap

Global
Variables

Code

Per-Thread State

Thread Control Book
(TCB)

Stack Information

Saved Registers

Thread Metadata

Stack



Execution Stack Example

```
A(int tmp) {  
    if (tmp<2)  
        B();  
    printf(tmp);  
}  
  
B() {  
    C(2);  
}  
  
C() {  
    A(2);  
}  
  
A(1);
```

Stack
Pointer →

A: tmp = 1
ret = exit

B: ret = A+2

C: ret = B+1

A: tmp = 2
ret = C+1

- Stack holds temporary results
- Permits recursive execution
- Crucial to modern languages



MIPS: SW Convention for Registers

- Before calling procedure

- Save caller-saves regs
- Save v0, v1
- Save ra

- After return, assume

- Callee-save reg OK
- gp, sp, fp restored
- Other things trashed

0	zero	constant 0	16	s0	callee saves
1	at	reserved for assembler	...		(callee must save)
2	v0	expression evaluation &	23	s7	
3	v1	function results	24	t8	temporary (cont'd)
4	a0	arguments	25	t9	
5	a1		26	k0	reserved for OS kernel
6	a2		27	k1	
7	a3		28	gp	Pointer to global area
8	t0	temporary: caller saves	29	sp	Stack pointer
...		(callee can clobber)	30	fp	frame pointer
15	t7		31	ra	Return Address (HW)



Motivational Example for Threads

- Imagine the following C program:

```
main() {  
    ComputePI("pi.txt");  
    PrintClassList("classlist.txt");  
}
```

- What is the behavior here?
 - Program would never print out class list
 - Why? ComputePI would never finish



User vs. Kernel Threads

- User threads: user-level threads library
 - 3 primary thread libraries
 - POSIX Pthreads
 - Windows threads
 - Java threads
- Kernel threads: supported by Kernel
 - Examples: Windows, Linux, Mac OS X



Kernel Threads

- Native threads supported directly by kernel
- Every thread can run or block independently
- One process may have several threads waiting on different things
- Downside: Expensive



User Threads

- Supported above the kernel, via a set of library calls at the user level
 - May have several user threads per kernel thread
 - May be scheduled non-preemptively relative to each other (only switch on **yield()**)
- Advantages: Cheap, Fast
- Downside: If kernel is single threaded, system call from any thread can block the entire task




Multithreading Models

- Many-to-One
- One-to-One
- Many-to-Many




Many-to-One

- User-level threads mapped to single kernel thread
- One thread blocking causes all to block
- Multithreads may not run in parallel on multicore system



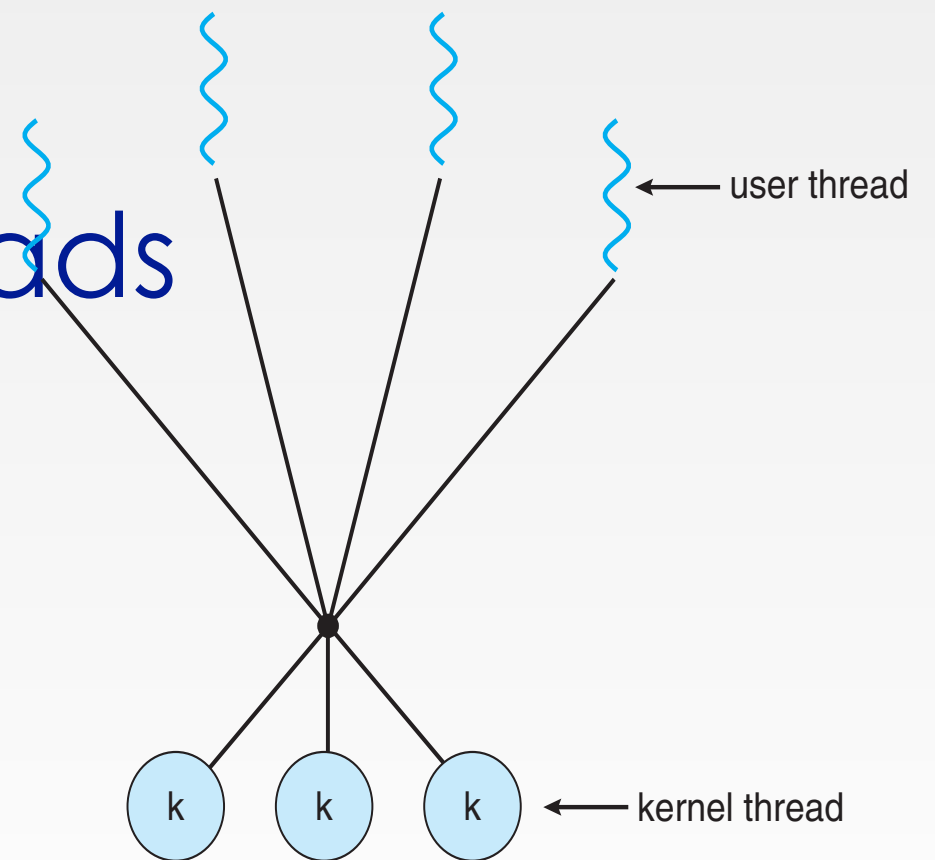
One-to-One

- Each user-level thread maps to kernel thread
- Creating a user-level thread creates a kernel thread
- More concurrency
- Overhead




Many-to-Many

- Allows many user-level threads to be mapped to many kernel threads
- Allows the OS to create a sufficient # of kernel threads



Two-Level Model

- Similar to *M-t-M*
- Allows a user thread to be bound to kernel thread



Thread Libraries

- Provides programmer with API for creating and managing threads
- Two primary ways of implementing
 - Library entirely in user space
 - Kernel-level library supported by the OS



Pthreads

- May be provided either as user/kernel level
- a POSIX standard API for thread creation and synchronization
- Specification, NOT implementation
- API specifies behavior of the thread library



Pthreads Example

```
#include <pthread.h>
#include <stdio.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    if (argc != 2) {
        fprintf(stderr, "usage: a.out <integer value>\n");
        return -1;
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr, "%d must be >= 0\n", atoi(argv[1]));
        return -1;
    }
}
```

```
/* get the default attributes */
pthread_attr_t attr;
/* create the thread */
pthread_create(&tid, &attr, runner, argv[1]);
/* wait for the thread to exit */
pthread_join(tid, NULL);

printf("sum = %d\n", sum);
}

/* The thread will begin control in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```



Pthread Example 2

- Joining 10 threads

```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```



Thread Pools

- Create a number of threads in a pool where they await work
- Advantages:
 - Slightly faster to service a request with an existing thread than create a new one
 - Allows the number of threads in the apps to be bound to the size of the pool
 - Allows different strategies for running task



OpenMP

- Set of compiler directives and an API for C/C++/Fortran
- Provides support for parallel programming in shared-memory environments
- Identifies parallel regions



OpenMP Example

```
void simple(int n, float *a, float *b)
{
    int i;

    #pragma omp parallel for
    for (i=1; i<n; i++) /* i is private by default */
        b[i] = (a[i] + a[i-1]) / 2.0;
}
```

create as many threads
as there are cores



Use of Threads

- Version of program with Threads (loose syntax)

```
main() {  
    ThreadFork(ComputePI, "pi.txt");  
    ThreadFork(PrintClassList, "classlist.txt");  
}
```

- What does ThreadFork() do?
 - Start independent thread running given procedure
- What is the behavior here?
 - As if there are two separate CPUs



Use of Threads

- Version of program with Threads (loose syntax)

```
main() {  
    ThreadFork (ComputePI, "pi.txt");  
    ThreadFork (PrintClassList, "classlist.txt");  
}
```

- What does ThreadFork() do?

– Sequence of execution



Time →

- What

– As if there are two separate CPUs



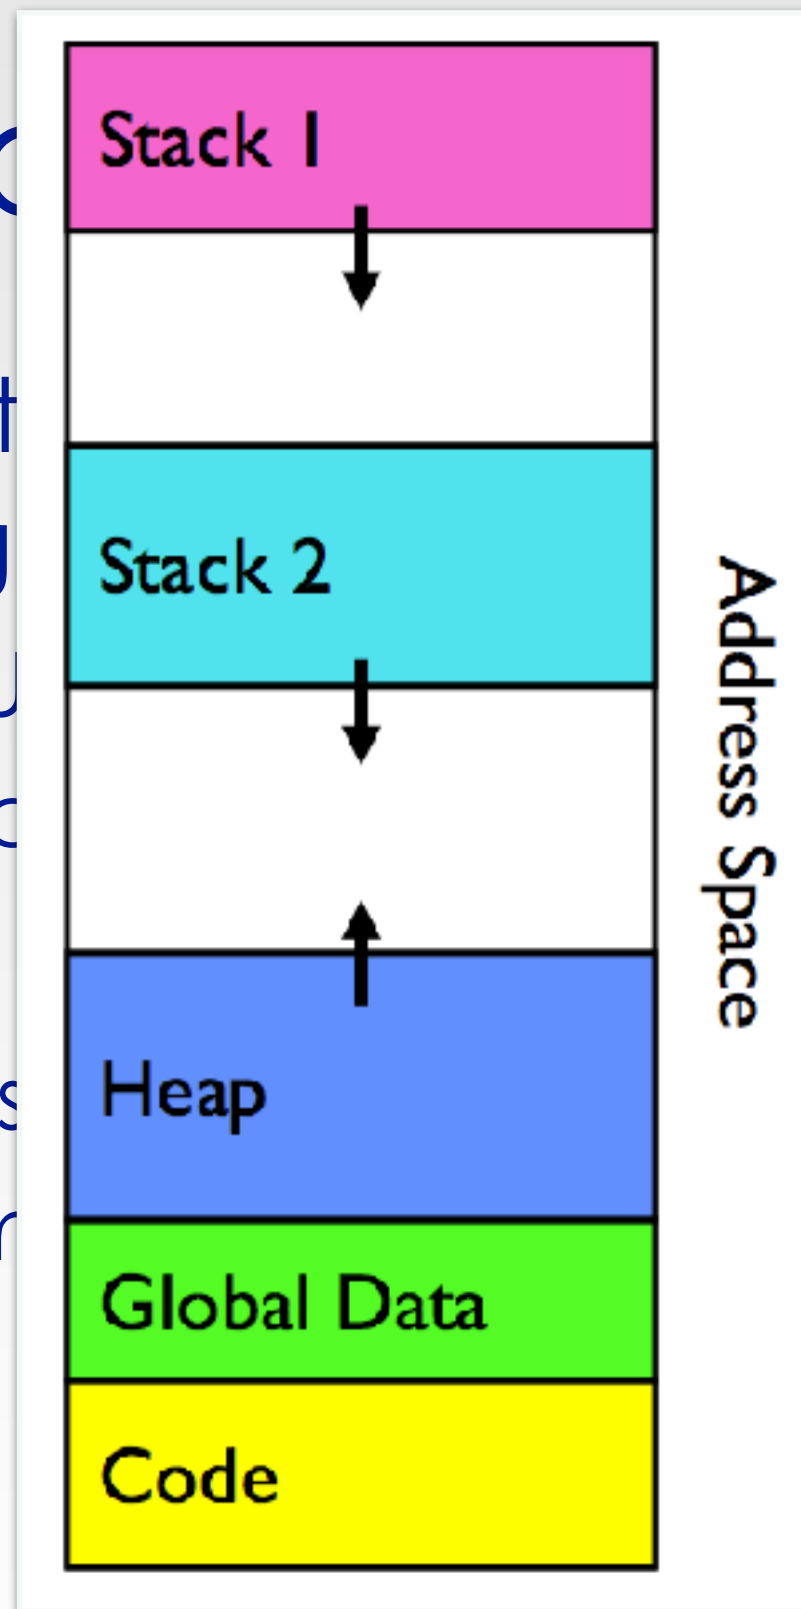
Memory Footprint: Two-Threads

- If we stopped this program and examined it with a debugger, we would see:
 - Two sets of CPU registers
 - Two sets of Stacks
- Questions:
 - How do we position stacks relative to each other?
 - What maximum size should we choose for the stacks
 - What happens if threads violate this?
 - How might you catch violations?



Memory For

- If we stopped the program with a debugger
 - Two sets of CPU registers
 - Two sets of Stack
- Questions:
 - How do we position the stacks?
 - What maximum size can we have?
 - What happens if they overlap?
 - How might you detect this?



Two-Threads

and examined it
e:

to each other?
oose for the stacks
his?

Actual Thread Operations

- `thread_fork(func, args)`
 - Create a new thread to run `func(args)`
- `thread_yield()`
 - Relinquish processor voluntarily
- `thread_join(thread)`
 - In parent, wait for forked thread to exit, then return
- `thread_exit`
 - Quit thread and clean up, wake up joiner if any
- `pThreads`
 - POSIX standard for thread programming



Dispatch Loop

- Conceptually, the dispatching loop of the operating system looks as follows:

```
Loop {  
    RunThread() ;  
    ChooseNextThread() ;  
    SaveStateOfCPU(curTCB) ;  
    LoadStateOfCPU(newTCB) ;  
}
```

- This is an infinite loop
 - Once could argue that this is all that the OS does
- Should we ever exit this loop?
 - When would that be?



Running a Thread

- Consider first position: `RunThread()`
- How do I run a thread?
 - Load its state (registers, PC, SP) into CPU
 - Load environment (virtual memory space, etc)
 - Jump to the PC
- How does the dispatcher get control back?
 - Internal events: thread returns control voluntarily
 - External events: thread gets preempted



Internal Events

- Blocking on I/O
- Waiting on a “signal” from other thread
 - Thread asks to wait and thus yields the CPU
- Thread executes a `yield()`
 - Thread volunteers to give up CPU

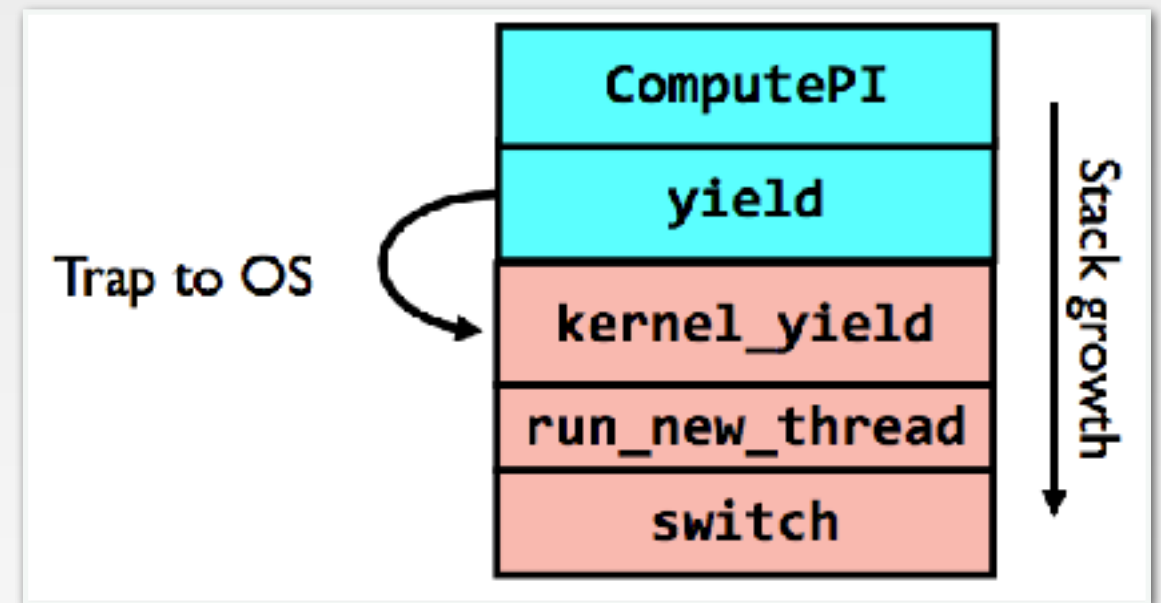
```
computePI() {  
    while (TRUE) {  
        ComputeNextDigit();  
        yield();  
    }  
}
```



Stack for Yielding Thread

- How do we run a new thread?

```
run_new_thread() {  
    newThread = PickNewThread();  
    switch(curThread, new Thread);  
    ThreadHouseKeeping();  
}
```



- How does dispatcher switch to a new thread?

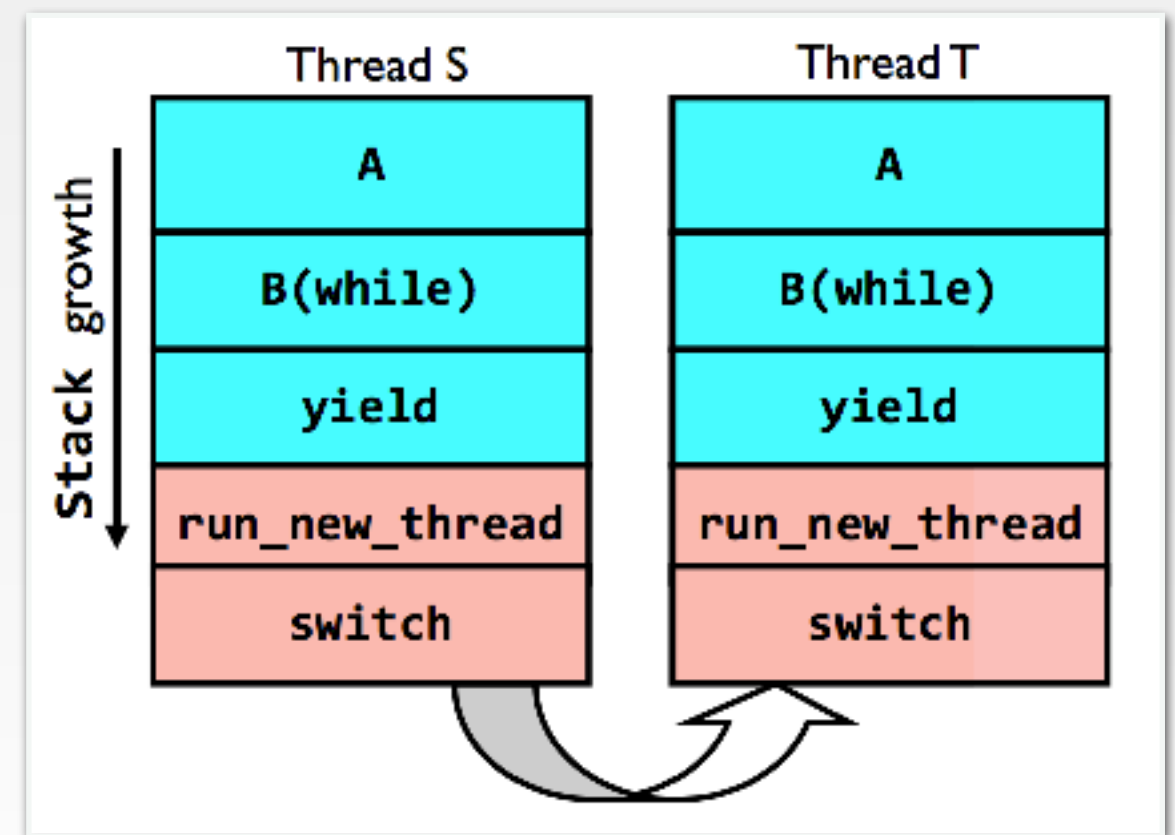
- Save PC, regs, SP
- Maintain isolation

What Do the Stacks Look Like?

- Consider the following code blocks:

```
proc (A) {  
    B();  
}  
  
proc (B) {  
    while (TRUE) {  
        yield();  
    }  
}
```

- Suppose we have two threads:
 - Threads S & T



Saving/Restoring State

- Often called “Context Switch”

```
switch(tCur, tNew){  
    /*unload old thread*/  
    TCB[tCur].regs.r7 = CPU.r7;  
    ...  
    TCB[tCur].regs.r0 = CPU.r0;  
    TCB[tCur].regs.sp = CPU.sp;  
    TCB[tCur].regs.retpc = CPU.retpc; /*return addr*/  
  
    /*load and execute new thread*/  
    CPU.r7 = TCB[tNew].regs.r7;  
    ...  
    CPU.r0 = TCB[tNew].regs.r0;  
    CPU.sp = TCB[tNew].regs.sp;  
    CPU.retpc = TCB[tNew].regs.retpc;  
    return; /*return to CPU.retpc*/  
}
```



Switch Details (cont.)

- What if you make a mistake in implementing switch?
 - Suppose you forget to save/restore register 32
 - Get intermittent failures depending on when context switch occurred and whether new thread uses register 32
 - System will give wrong result without warning
- Can you devise an exhaustive test to test switch code?
 - No.
 - Too many combinations and inter-leavings



Switch Details (cont.)

- Cautionary tale:
 - For speed, Topaz kernel saved one instruction in **switch{ }**
 - Carefully documented. Only works as long as kernel size < 1MB
 - What happens:
 - Time passed, people forget
 - Later, they added features to kernel (no one removes features)
 - Very weird behavior started happening



Some Numbers

- Frequency of performing context switch:
 - 10-100ms
- Context switch time in Linux: 3-4 μ secs
 - Thread switching faster than process switching (100 ns)
 - But switching across cores ~2X more expensive than within-core
- Context switch time increases sharply with size of working set
 - Can increase 100x or more



Some Numbers (cont.)

- Moral: context switching depends mostly on cache limits and the process or thread's hunger for memory
- Many processes are multi-threaded, so thread context switches may be either **within-process** or **across-processes**



Some Numbers (cont.)

- More on c threads mostly
- Many threads, so
- with ner

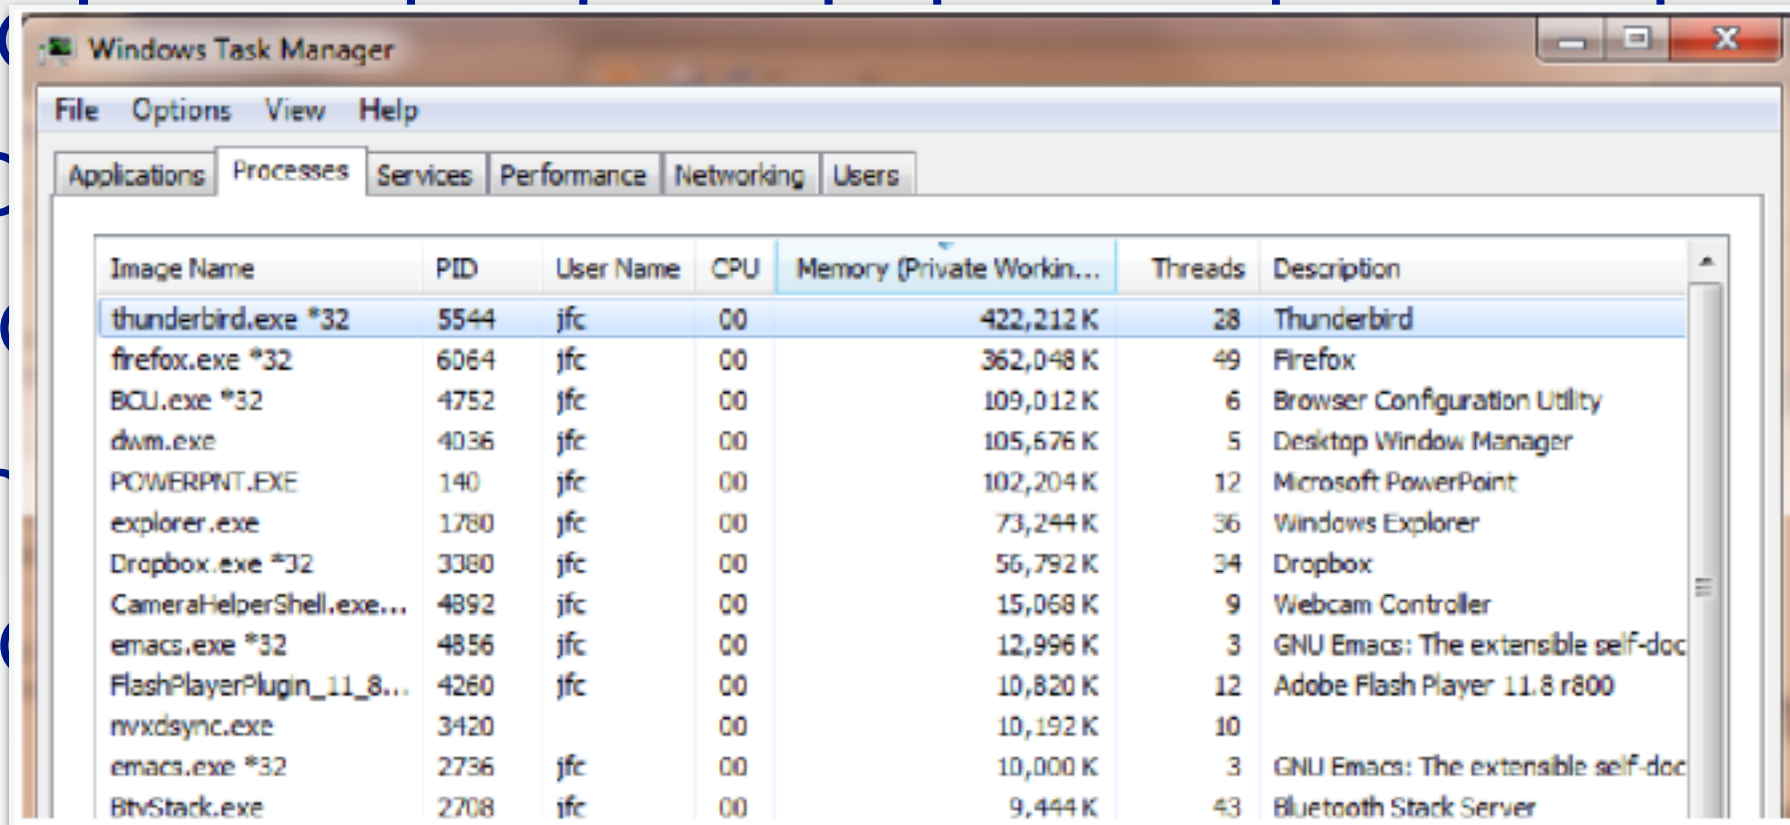
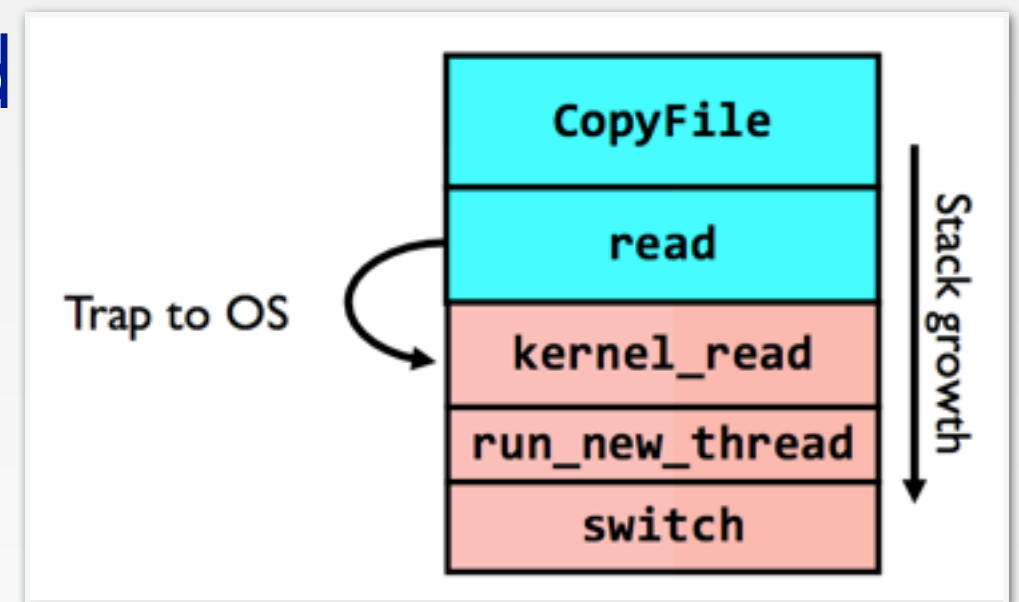


Image Name	PID	User Name	CPU	Memory (Private Workin...	Threads	Description
thunderbird.exe *32	5544	jfc	00	422,212 K	28	Thunderbird
firefox.exe *32	6064	jfc	00	362,048 K	49	Firefox
BCU.exe *32	4752	jfc	00	109,012 K	6	Browser Configuration Utility
dwm.exe	4036	jfc	00	105,676 K	5	Desktop Window Manager
POWERPNT.EXE	140	jfc	00	102,204 K	12	Microsoft PowerPoint
explorer.exe	1780	jfc	00	73,244 K	36	Windows Explorer
Dropbox.exe *32	3380	jfc	00	56,792 K	34	Dropbox
CameraHelperShell.exe...	4892	jfc	00	15,068 K	9	Webcam Controller
emacs.exe *32	4856	jfc	00	12,996 K	3	GNU Emacs: The extensible self-doc
FlashPlayerPlugin_11_8...	4260	jfc	00	10,820 K	12	Adobe Flash Player 11.8 r800
nvxdsync.exe	3420		00	10,192 K	10	
emacs.exe *32	2736	jfc	00	10,000 K	3	GNU Emacs: The extensible self-doc
BtvStack.exe	2708	jfc	00	9,444 K	43	Bluetooth Stack Server

What happens when thread blocks on I/O?

- What happens when a thread requests a block of data from the FS?
 - User code invokes a system call
 - Read operation is initiated
 - Run new thread/switch
- Thread communication
 - Similar to Networking
 - or Wait for Signal/Join



External Events

- What happens if thread never does any I/O, never waits, or never yield control?
 - Could the **ComputePI** program grab all resources and never release processor?
 - Must find ways that dispatcher can regain control.
- Answer: Utilize External Events



External Events

- Interrupts: signals from HW or SW that stop the running code and jump to kernel
- Timer: like an alarm clock that goes off every some may milliseconds



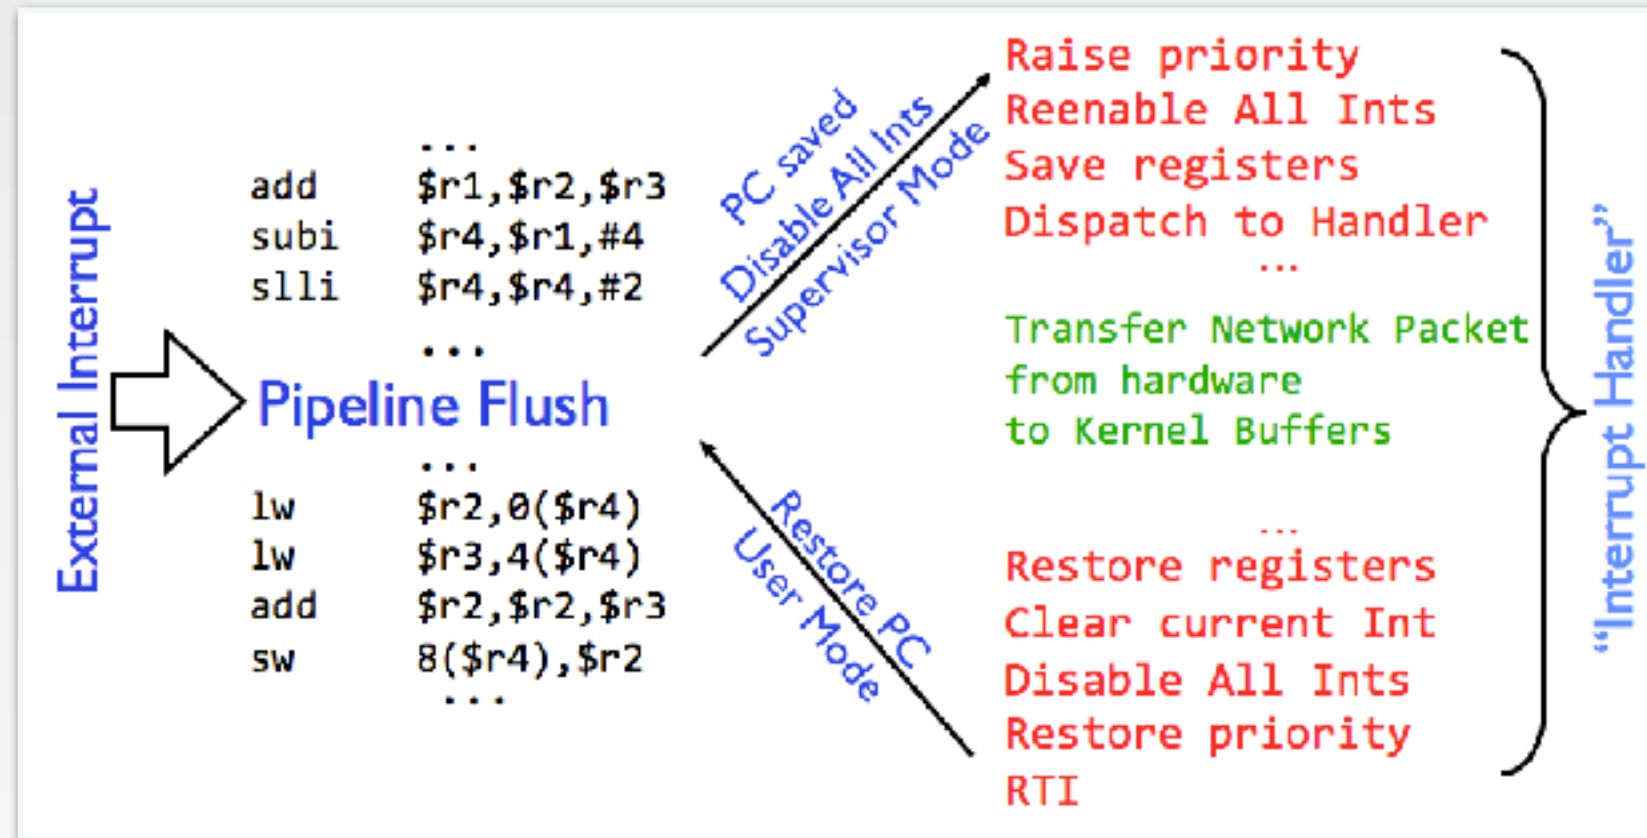
External Events

- Interrupts: signals from HW or SW that stop the running code and jump to kernel
- Timer: like an alarm clock that goes off every some may milliseconds

If we make sure that external events occur frequently enough, can ensure dispatcher runs



Example: Network Interrupt



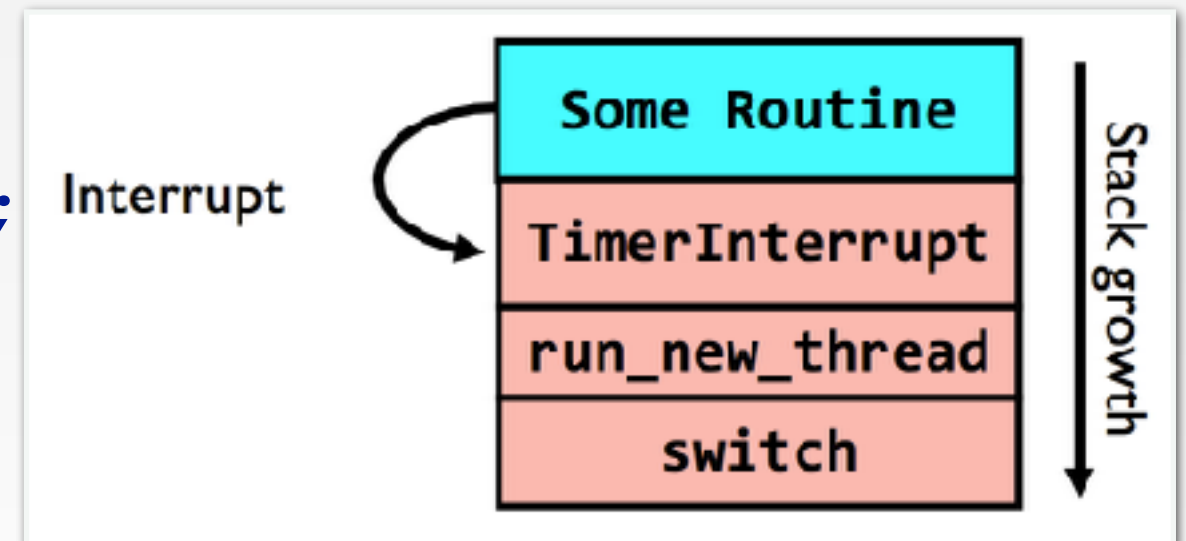
- An interrupt is a HW-invoked context switch
 - No separate step to choose what to run next
 - Always run the interrupt handler immediately



Example: Use Timer to Return Control

- Solution to our dispatcher problem
 - Use the timer interrupt to force scheduling decisions
- Timer interrupt routine:

```
TimerInterrupt() {  
    DoPeriodicHouseKeeping();  
    run_new_thread();  
}
```



Thread Abstraction



- Illusion:

Infinite number of processors

- Reality:

Threads execute w/ variable speed

- Programs must be designed to work with any schedule



Programmer vs. Processor View

Programmer's View	Possible Execution #1
.	.
.	.
.	.
$x = x + 1;$	$x = x + 1;$
$y = y + x;$	$y = y + x;$
$z = x + 5y;$	$z = x + 5y;$
.	.
.	.
.	.



Programmer vs. Processor View


Programmer's View	Possible Execution #1	Possible Execution #2
.	.	.
.	.	.
.	.	.
$x = x + 1;$	$x = x + 1;$	$x = x + 1$
$y = y + x;$	$y = y + x;$
$z = x + 5y;$	$z = x + 5y;$	thread is suspended
.	.	other thread(s) run
.	.	thread is resumed
.
		$y = y + x$
		$z = x + 5y$




Programmer vs. Processor View

Programmer's View	Possible Execution #1	Possible Execution #2	Possible Execution #3
.	.	.	.
.	.	.	.
.	.	.	.
$x = x + 1;$	$x = x + 1;$	$x = x + 1$	$x = x + 1$
$y = y + x;$	$y = y + x;$	$y = y + x$
$z = x + 5y;$	$z = x + 5y;$	thread is suspended
.	.	other thread(s) run	thread is suspended
.	.	thread is resumed	other thread(s) run
.	thread is resumed
		$y = y + x$
		$z = x + 5y$	$z = x + 5y$


Possible Executions



a) One execution

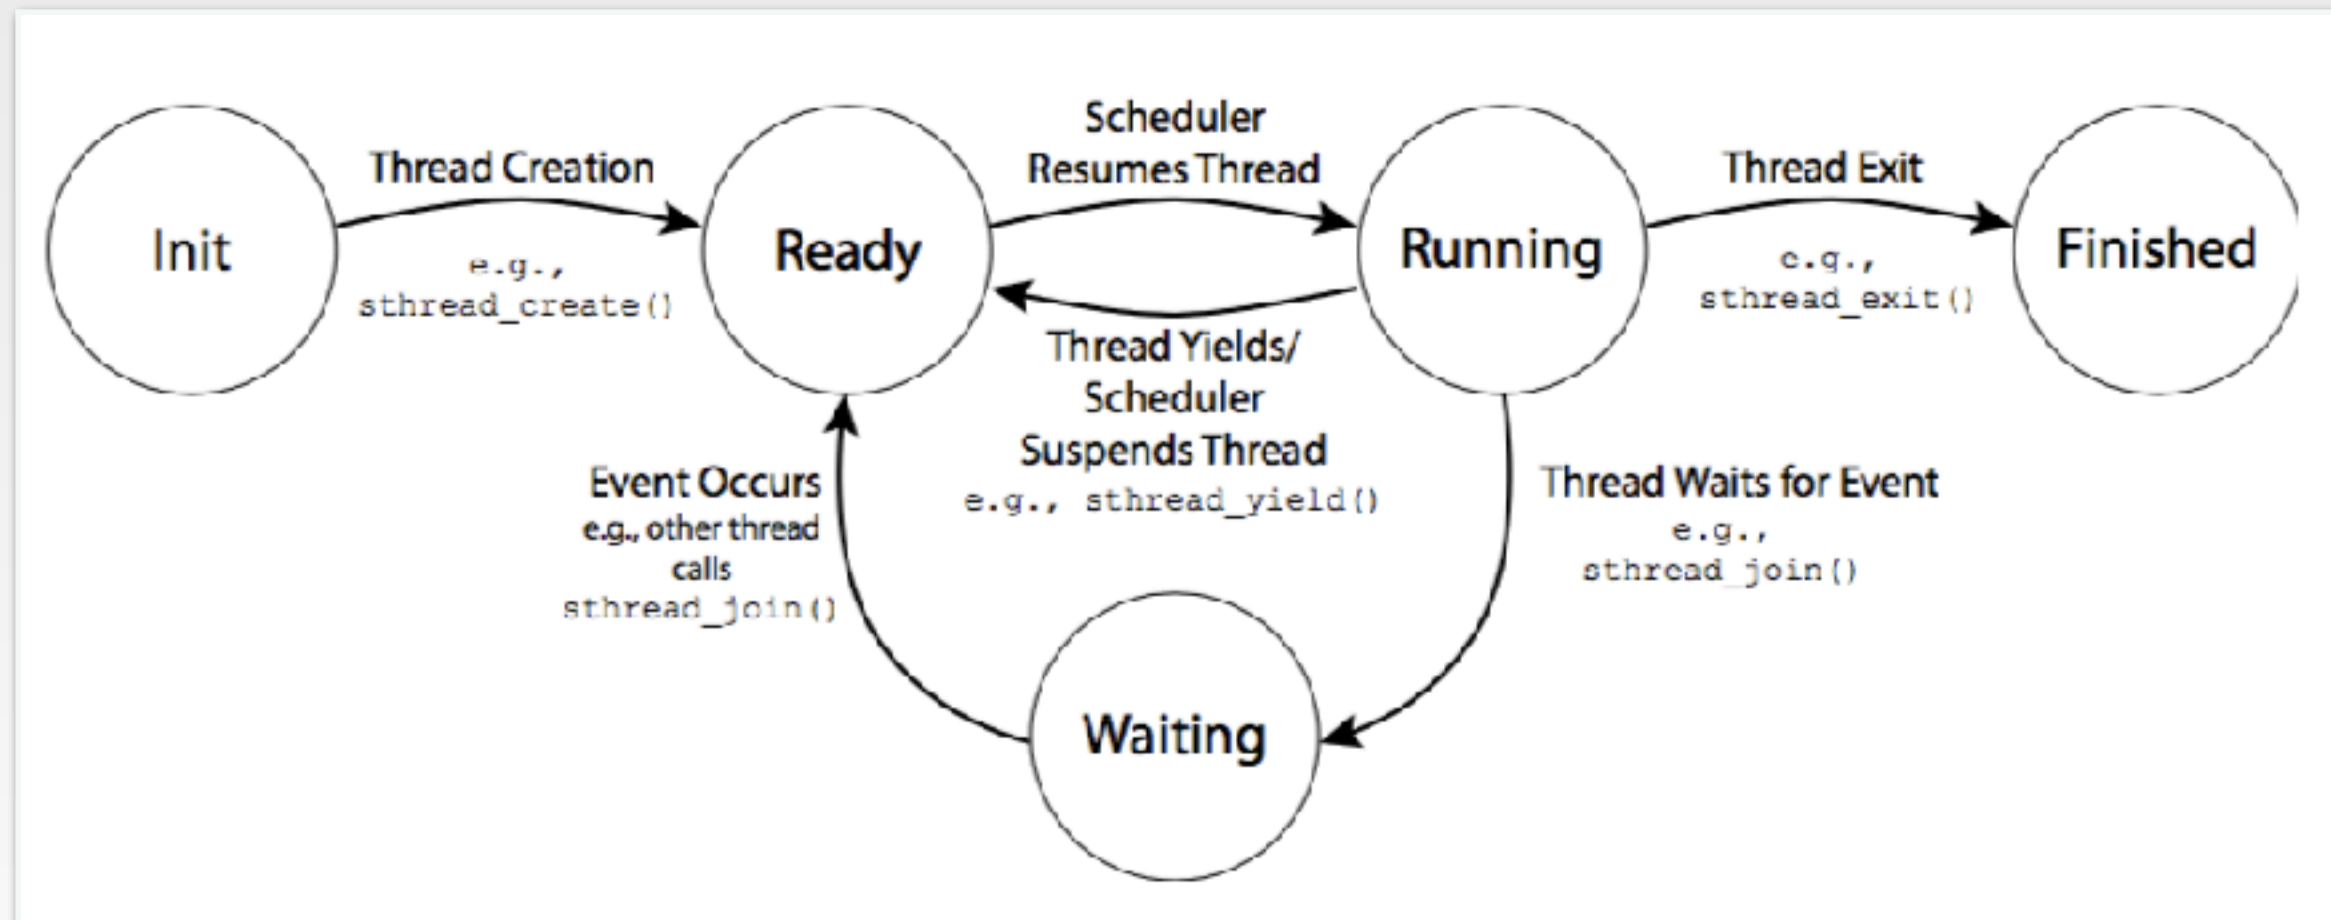


b) Another execution



c) Another execution

Thread Lifecycle



Summary

- Processes have two parts
 - One or more Threads (concurrency)
 - Address Spaces (protection)
- Concurrency accomplished by multiplexing CPU Time:
 - Unloading current thread (PC, registers)
 - Loading new thread (PC, registers)
 - Such context switching may be voluntary or involuntary



Discussion

- Multi-processing/programming/threading
 - Definition:
 - Multiprocessing → multiple CPUs
 - Multiprogramming → Multiple Jobs or Processes
 - Multithreading → Multiple Threads per Process



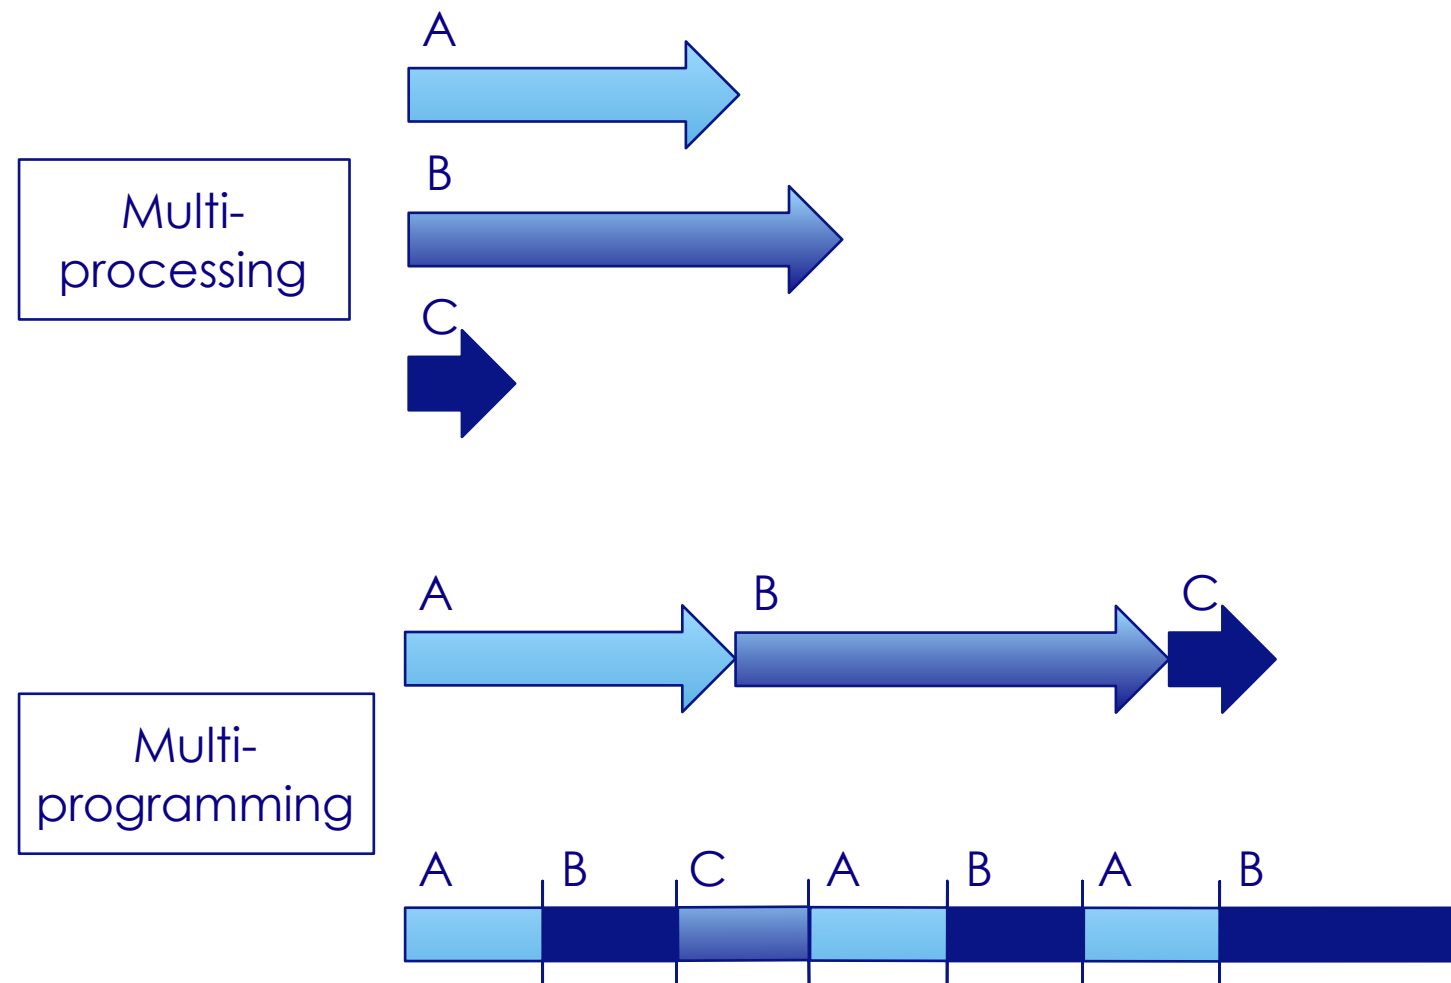
Discussion (cont.)

- What does it mean to run threads “concurrently”?
 - Scheduler is free to run threads in any order and interleaving: FIFO, Random, etc.
 - Dispatcher can choose to run each thread to completion or time-slice in big chunks or small chunks



Discussion (cont.)

- Who
- “con
- Sch
- and
- Dis
- con
- chu



order
lead to
or small