



Operating Systems

Dr. Shu Yin

Part II: Process Management

- Processes
- Threads
- Process Synchronization
- CPU Scheduling
- Deadlocks



Goals

- Interprocess Communication
- Process Synchronization



Interprocess Communication

- Process within a system may be *independent* or *cooperating*
- Cooperating process can affect or be affected by other processes, including sharing data



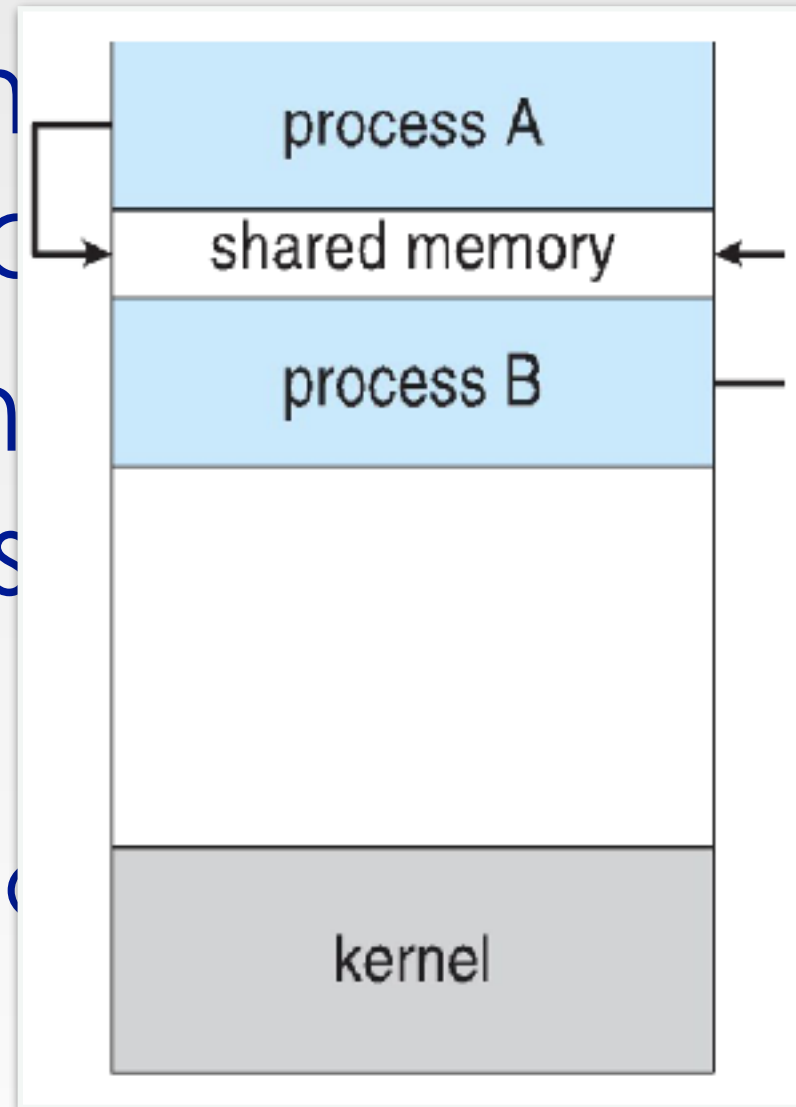
Interprocess Communication (cont.)

- Reasons for cooperating processes:
 - Information sharing
 - Computation speedup
 - Modularity
 - Convenience (e.g. editing, printing, compiling)
- Cooperating processes need IPC
- Two models of IPC
 - Shared memory
 - Message passing



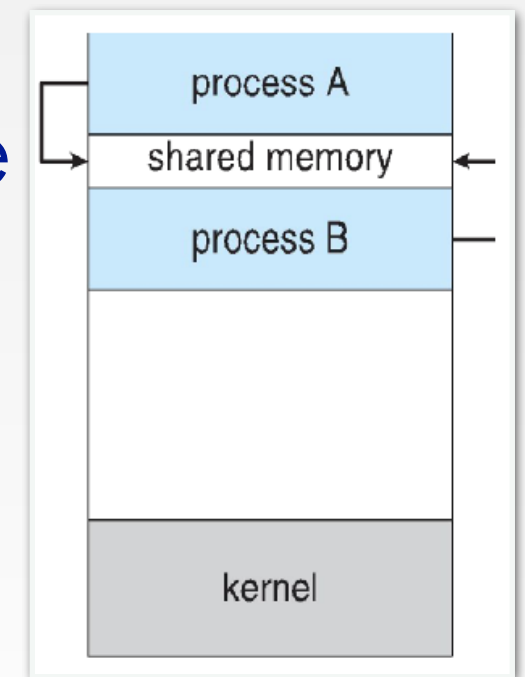
IPC: Shared Memory

- An area of memory that is shared among the processes that communicate
- The communication is done under the control of the kernel
- Major issue:
 - Synchronization (to be discussed later)



IPC: Shared Memory (cont.)

- Producer-Consumer Problem
 - unbounded-buffer
 - places no practical limit on the size of the buffer
 - bounded-buffer
 - assumes that there is a fixed buffer size



IPC: Shared Memory (cont.)

- Bounded-buffer
- Shared-Memory Solution

```
#define BUFFER_SIZE 10
typedef struct{
...
}item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

shared data

```
item next_consumed;
while (true){
    while (in == out)
        /*do nothing*/
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    /*consume the item in the next
    consumed*/
}
```

consumer

```
item next_produced;
while (true){
    /*produce an item in next produced*/
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /*do nothing*/
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```

producer



IPC: Shared Memory (cont.)

- Bounder-Buffer
- How many elements in the buffer can be used at most a a given time?

```
item next_produced;
while (true){
    /*produce an item in next produced*/
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /*do nothing*/
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```

producer

```
item next_consumed;
while (true){
    while (in == out)
        ;/*do nothing*/
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    /*consume the item in the next
    consumed*/
}
```

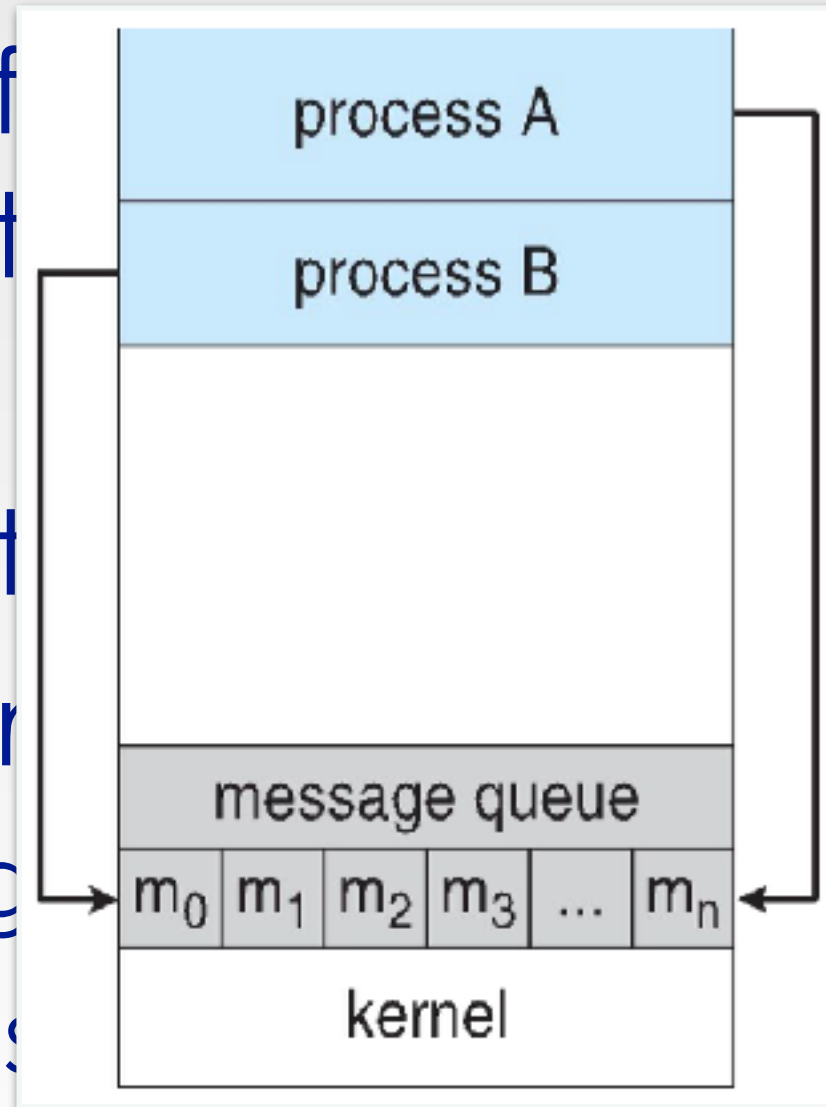
consumer



上海科技大学
ShanghaiTech University

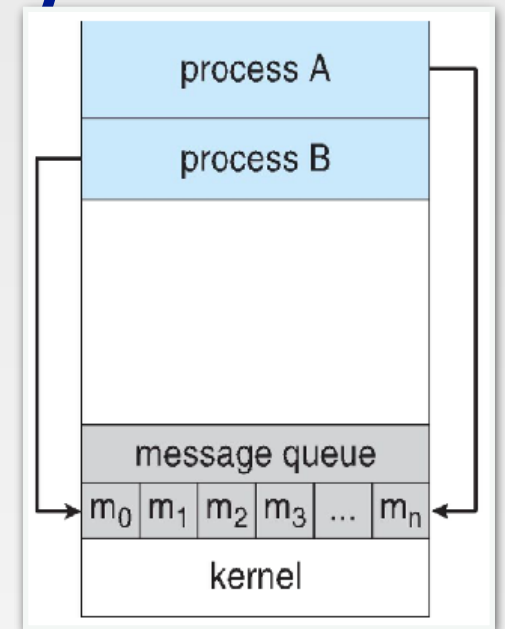
IPC: Message Passing

- Mechanism for communication and synchronization
- Message system
- IPC facility provides two operations:
 - **send**(message)
 - **receive**(message)
- Message size: fixed or variable



IPC: Message Passing (cont.)

- If processes A and B wish to communicate
 - Establish a communication link
 - Exchange messages via send/receive
- Implementation issues:
 - How are links established
 - Can a link be associated with one or more processes
 - How many links What's the capacity of a link
 - Is the size of message fixed/variable
 - Is a link unidirectional or bi-directional



IPC: Message Passing (cont.)

- Implementation of communication link
 - Physical
 - Shared memory
 - HW bus
 - Network

Before we further discuss the MP, the communication should be discussed first



Direct Communication

- Processes must name each other explicitly
 - **send**(A, message)
 - Send a message to process A
 - **receive**(B, message)
 - Receive a message from process B
- Properties of communication link
 - Established automatically
 - Exists exactly one link
 - May be uni-directional, but usually bi-directional
 - A link is associated with exactly one pair



Indirect Communication

- Messages are directed and received from mailboxes (also referred to as ports)
 - Each mailbox has a unique ID
 - Processes can communicate only if they share a mailbox
- Properties of communication link
 - Share a common mailbox
 - May be associated with many processes
 - Each pair may share several links
 - May be uni-directional or bi-directional



Indirect Communication (cont.)

- Operations
 - Create a new mailbox (port)
 - Send and Receive messages through mailbox
 - Destroy a mailbox
- Primitives are defined as
 - **send**(A, message)
 - Send a message to mailbox A
 - **receive**(A, message)
 - Receive a message from mailbox A



Indirect Communication (cont.)

- Mailbox sharing
 - P_1 , P_2 , and P_3 share mailbox A
 - P_1 sends; P_2 and P_3 receive
 - Who gets the message?
- Solutions
 - Allow a link to be associated with at most two processes
 - Allow only one process at a time to execute a receive operation
 - Allow the system to select arbitrarily the receiver
 - Sender is notified who the receiver was



Synchronization

- Message passing may be either blocking or non-blocking
- Blocking - synchronous
 - Blocking send
 - Blocking receive
- Non-blocking - asynchronous
 - Non-blocking send
 - Non-blocking receive: Valid/ Null
- Different combinations possible
 - Both S/R are blocking, then Rendezvous



Message Passing (cont.)

- Producer-consumer becomes trivial

```
item next_produced;  
while (true){  
    /*produce an item in next  
    produced*/  
    send (next_produced)  
}
```

producer

```
item next_consumed;  
while (true){  
    while (in == out)  
        receive(next_consumed);  
    /*consume the item in the  
    next consumed*/  
}
```

consumer



Buffering

- Queues of messages attached to the link
- Implemented in one of three ways
 - Zero capacity
 - Bounded capacity
 - Unbounded capacity



Example of IPC: POSIX

- POSIX shared memory

- Process first creates shared memory segment

```
shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);
```

- Also used to open an existing segment to share it

- Set the size of the object

```
ftruncate(shm_fd, 4096);
```

- Now the process could write to the shared memory

```
sprintf(shared_memory, "Writing to shared memory");
```



Example of IPC: POSIX

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* strings written to shared memory */
    const char *message_0 = "Hello";
    const char *message_1 = "World!";

    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* create the shared memory object */
    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

    /* configure the size of the shared memory object */
    ftruncate(shm_fd, SIZE);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

    /* write to the shared memory object */
    sprintf(ptr, "%s", message_0);
    ptr += strlen(message_0);
    sprintf(ptr, "%s", message_1);
    ptr += strlen(message_1);

    return 0;
}
```

Producer

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* open the shared memory object */
    shm_fd = shm_open(name, O_RDONLY, 0666);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

    /* read from the shared memory object */
    printf("%s", (char *)ptr);

    /* remove the shared memory object */
    shm_unlink(name);

    return 0;
}
```

Consumer



上海科技大学
ShanghaiTech University

Communications in C-S Systems

- Sockets
- Remote Procedure Calls
- Pipes



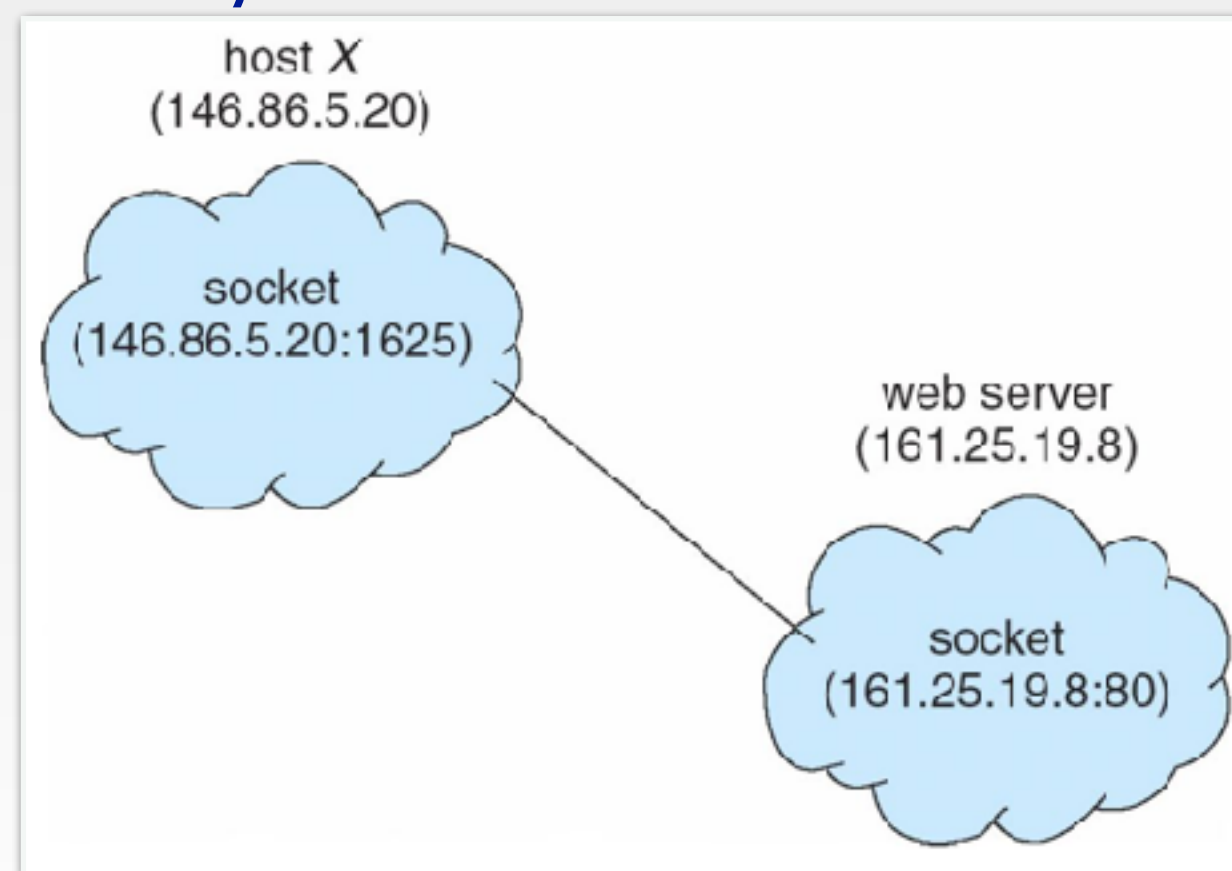
Sockets

- An endpoint for communication
- Concatenation of IP and Port
- The socket 172.16.254.1:22
 - port 22 on host 172.16.254.1
- Communication consists between a pair of sockets
- All ports below 1024 are well known
 - Used for standard services
- Special IP address 127.0.0.1
 - Loopback
 - Refers to system on which process is running



Socket Communication

- Three types
 - TCP (Connection-oriented)
 - UDP (Connectionless)
 - MulticastSocket



Connection-oriented Communication

- Session is established before transferring
- Delivered in the same order as it was sent
- Acknowledge after successful delivery
- TCP



Connectionless Communication

- Message sent from one end to another w/o prior arrangement
- Send w/o ensuring if available and ready
- IP, UDP
- Lower overhead
- Allows for multicast and broadcast
- Error correction to reduce error effects



Sockets in C

server

```
/* A simple server in the internet domain using TCP
   The port number is passed as an argument */
#include <stdio.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>

void error(char *msg)
{
    perror(msg);
    exit(1);
}

int main(int argc, char *argv[])
{
    int sockfd, newsockfd, portno, cliilen;
    char buffer[256];
    struct sockaddr_in serv_addr, cli_addr;
    int n;
    if (argc < 2) {
        fprintf(stderr, "ERROR, no port provided\n");
        exit(1);
    }
    sockfd = socket(AF_INET, SOCK_STREAM, 0);
    if (sockfd < 0)
        error("ERROR opening socket");
    bzero((char *) &serv_addr, sizeof(serv_addr));
    portno = atoi(argv[1]);
    serv_addr.sin_family = AF_INET;
    serv_addr.sin_addr.s_addr = INADDR_ANY;
    serv_addr.sin_port = htons(portno);
    if (bind(sockfd, (struct sockaddr *) &serv_addr,
        sizeof(serv_addr)) < 0)
        error("ERROR on binding");
    listen(sockfd, 5);
    cliilen = sizeof(cli_addr);
    newsockfd = accept(sockfd, (struct sockaddr *) &cli_addr, &cliilen);
    if (newsockfd < 0)
        error("ERROR on accept");
    bzero(buffer, 256);
    n = read(newsockfd, buffer, 255);
    if (n < 0) error("ERROR reading from socket");
    printf("Here is the message: %s\n", buffer);
    n = write(newsockfd, "I got your message", 18);
    if (n < 0) error("ERROR writing to socket");
    return 0;
}
```

client

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <netdb.h>

void error(char *msg)
{
    perror(msg);
    exit(0);
}

int main(int argc, char *argv[])
{
    int sockfd, portno, n;
    struct sockaddr_in serv_addr;
    struct hostent *server;

    char buffer[256];
    if (argc < 3) {
        fprintf(stderr, "usage %s hostname port\n", argv[0]);
        exit(0);
    }
    portno = atoi(argv[2]);
    sockfd = socket(AF_INET, SOCK_STREAM, 0);
    if (sockfd < 0)
        error("ERROR opening socket");
    server = gethostbyname(argv[1]);
    if (server == NULL) {
        fprintf(stderr, "ERROR, no such host\n");
        exit(0);
    }
    bzero((char *) &serv_addr, sizeof(serv_addr));
    serv_addr.sin_family = AF_INET;
    bcopy((char *) server->h_addr,
        (char *) &serv_addr.sin_addr.s_addr,
        server->h_length);
    serv_addr.sin_port = htons(portno);
    if (connect(sockfd, (struct sockaddr *) &serv_addr, sizeof(serv_addr)) < 0)
        error("ERROR connecting");
    printf("Please enter the message: ");
    bzero(buffer, 256);
    fgets(buffer, 255, stdin);
    n = write(sockfd, buffer, strlen(buffer));
    if (n < 0)
        error("ERROR writing to socket");
    bzero(buffer, 256);
    n = read(sockfd, buffer, 255);
    if (n < 0)
        error("ERROR reading from socket");
    printf("%s\n", buffer);
    return 0;
}
```



Remote Procedure Calls

- RPC abstracts procedure calls between processes on networked systems
 - Uses port for service differentiation
- Stubs
 - Client-side proxy for the actual procedure on the server



Remote Procedure Calls (cont.)

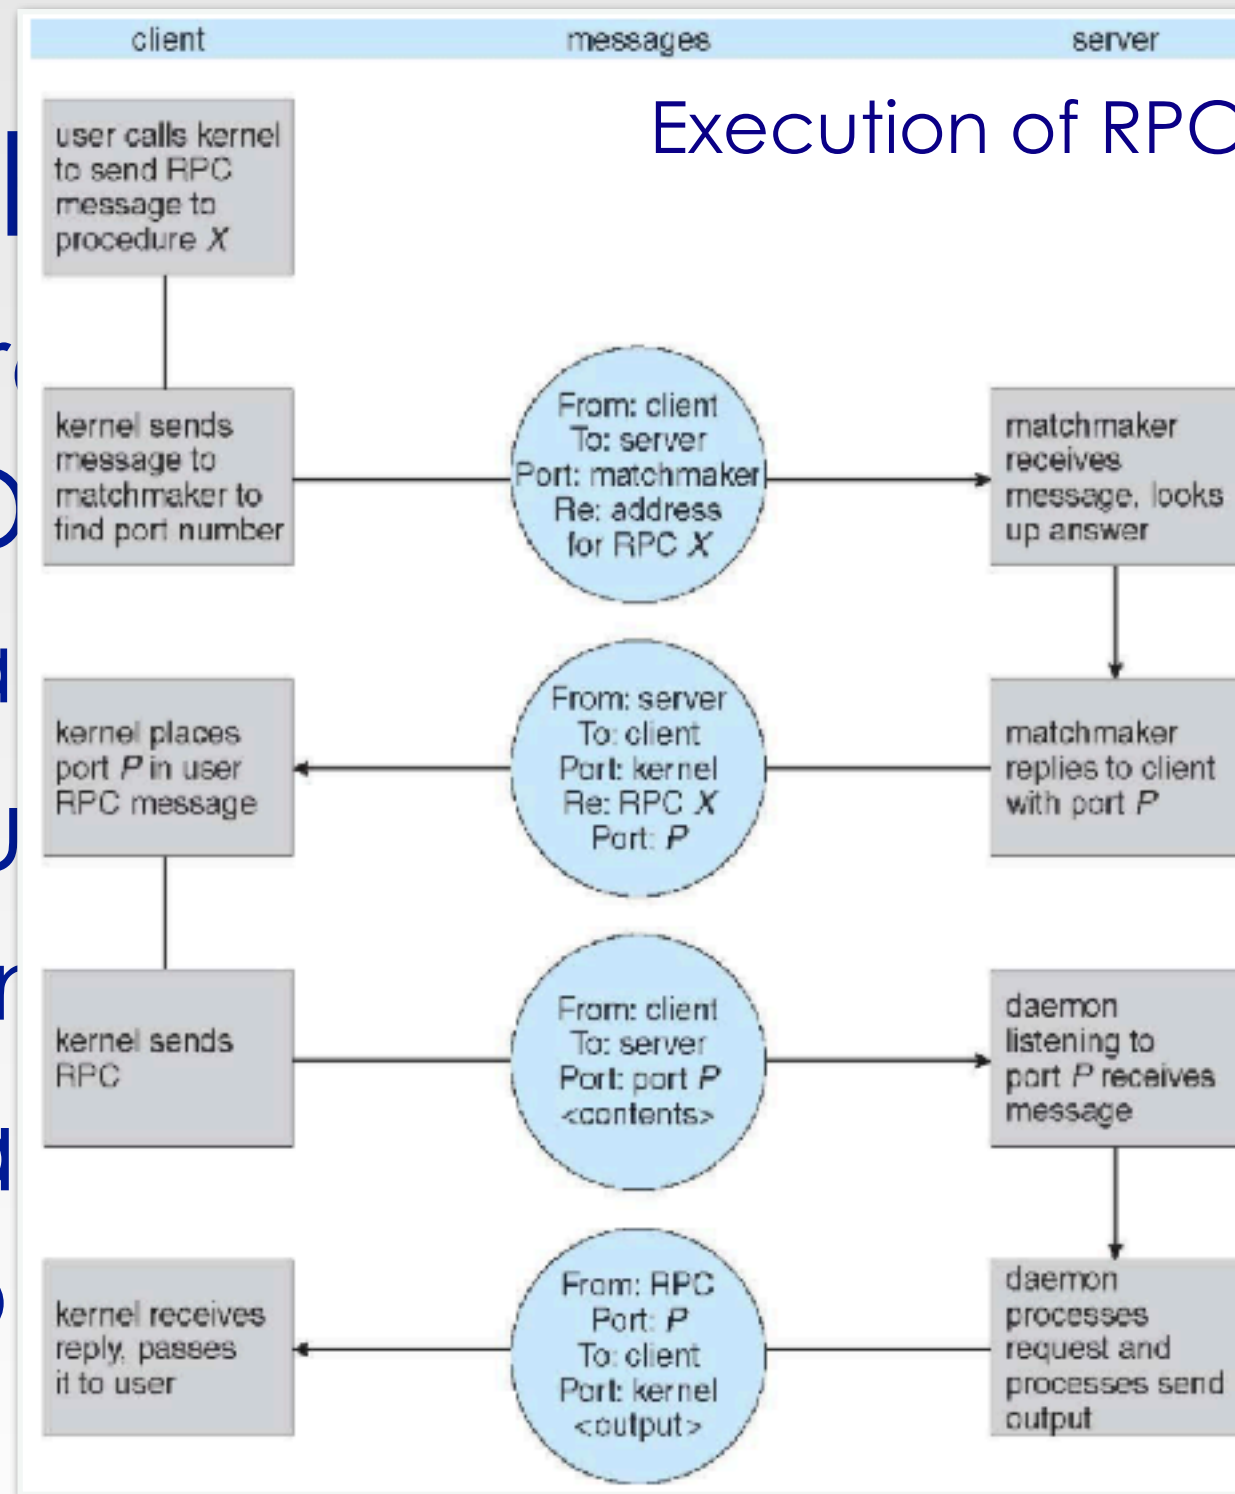
- Client-side stub locates the server and marshalls the parameters
- Server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server



Remote I/O

- Data representation (External Data Format)
 - Big-endian vs. Little-endian
- More failure modes (Remote vs. Local)
 - Exactly-once semantics
- OS typically provides service to

Execution of RPC



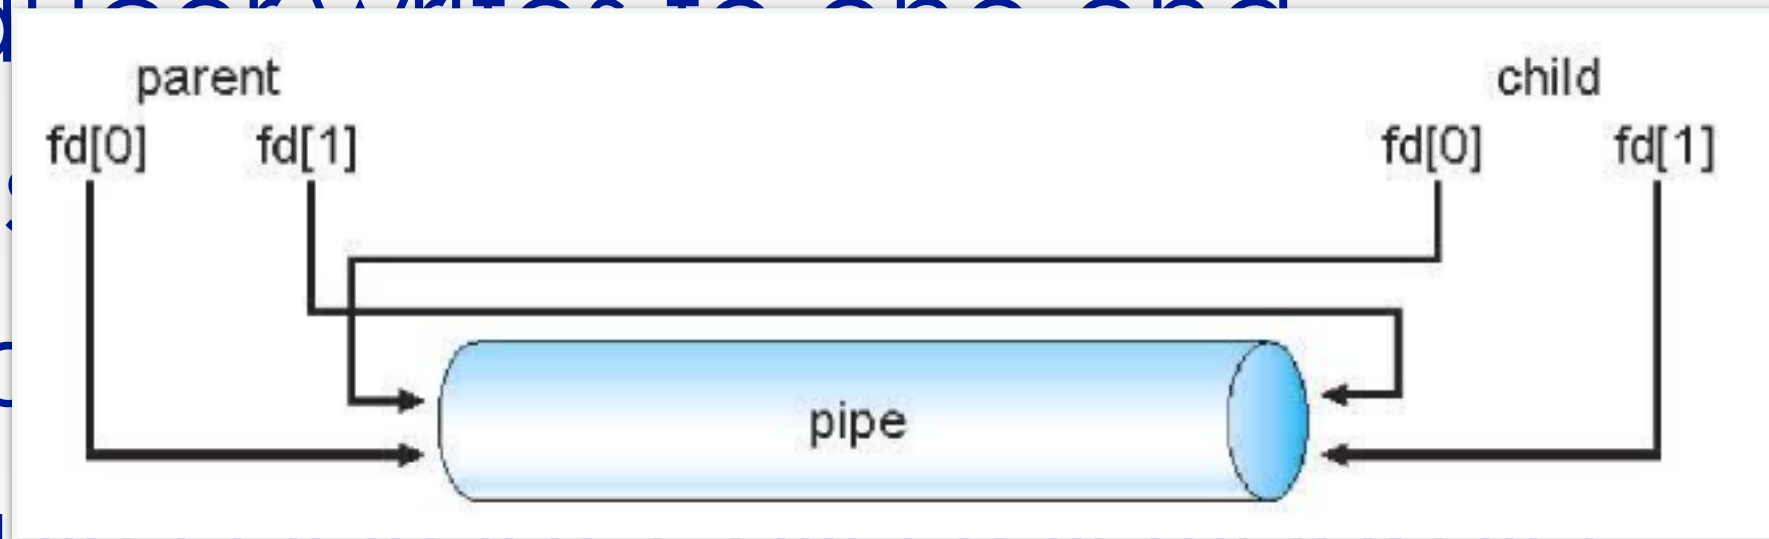
Pipes

- Acts as a conduit allowing two processes to communicate
- Ordinary pipes
 - Can NOT be accessed from the outside
- Named pipes
 - Can be accessed w/t a parent-child relationship



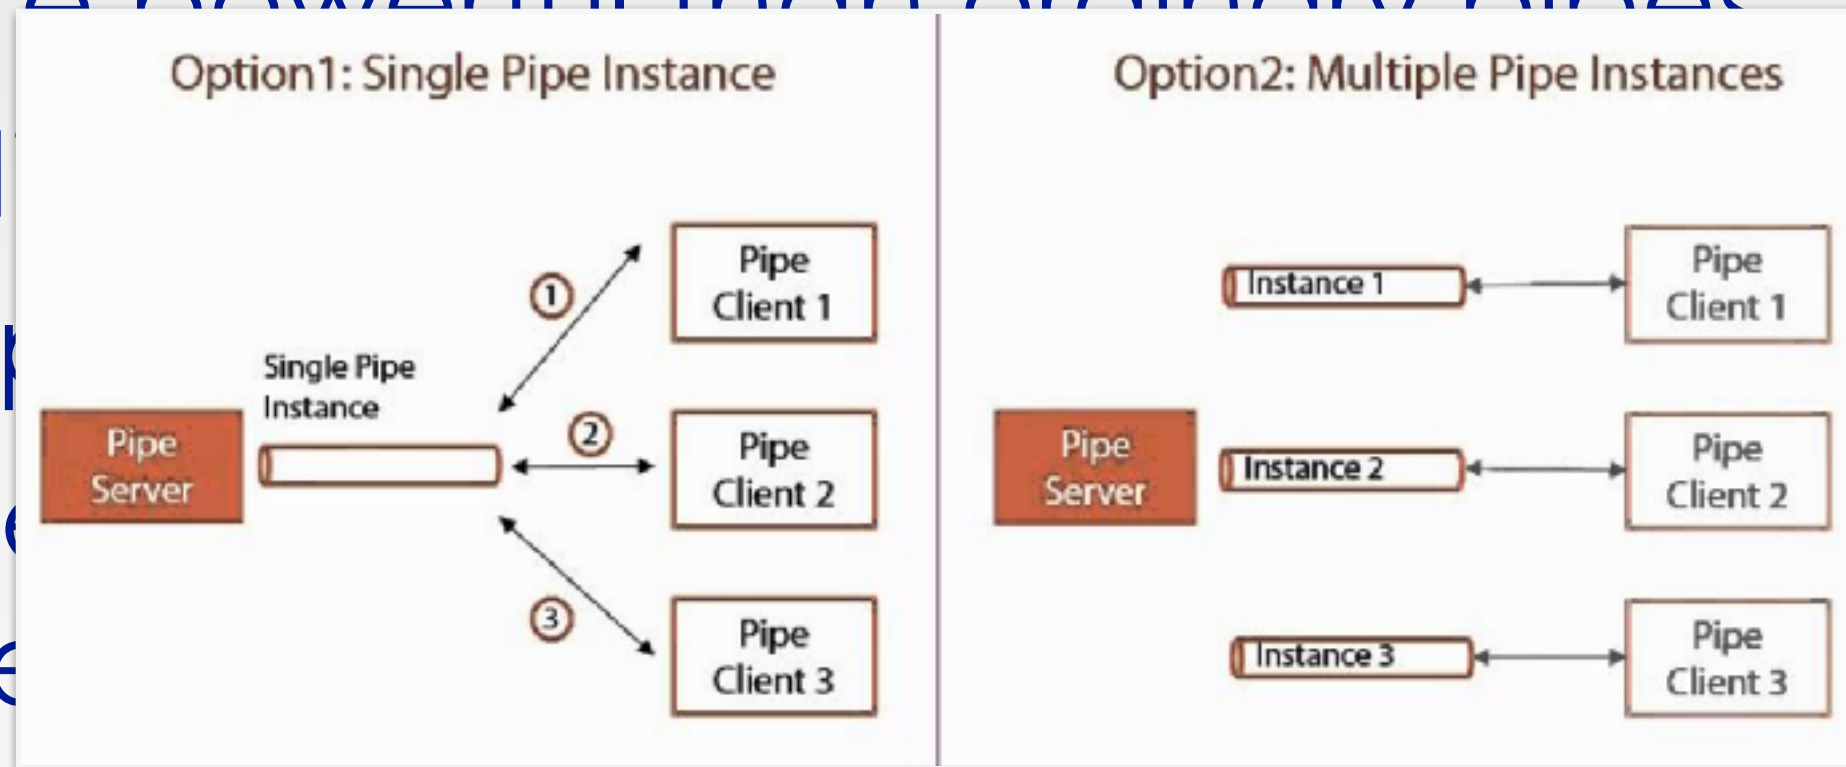
Ordinary Pipes

- Allow communication in standard Producer-Consumer style
- Producer writes to one end
- Consumer reads from the other
- Uni-directional
- Require parent-child relationship



Named Pipes

- More powerful than ordinary pipes
- Bi-directional
- No process ID limitation
- Several pipe instances
- Provided on both UNIX and Windows



Summary

- IPC
 - Shared Memory
 - Message Passing
 - Sockets

