



Operating Systems

Dr. Shu Yin

Part II: Process Management

- Processes
- Threads
- Process Synchronization
- CPU Scheduling
- Deadlocks



Goals

- Introduce the Critical Section Problem
- Both SW and HW Solutions of the C-S Problem
- Classical Problems of Synchronization
- Tools to Solve Process Sync. Problems



Recall:

Producer-Consumer Problem

- Paradigm

- Producer

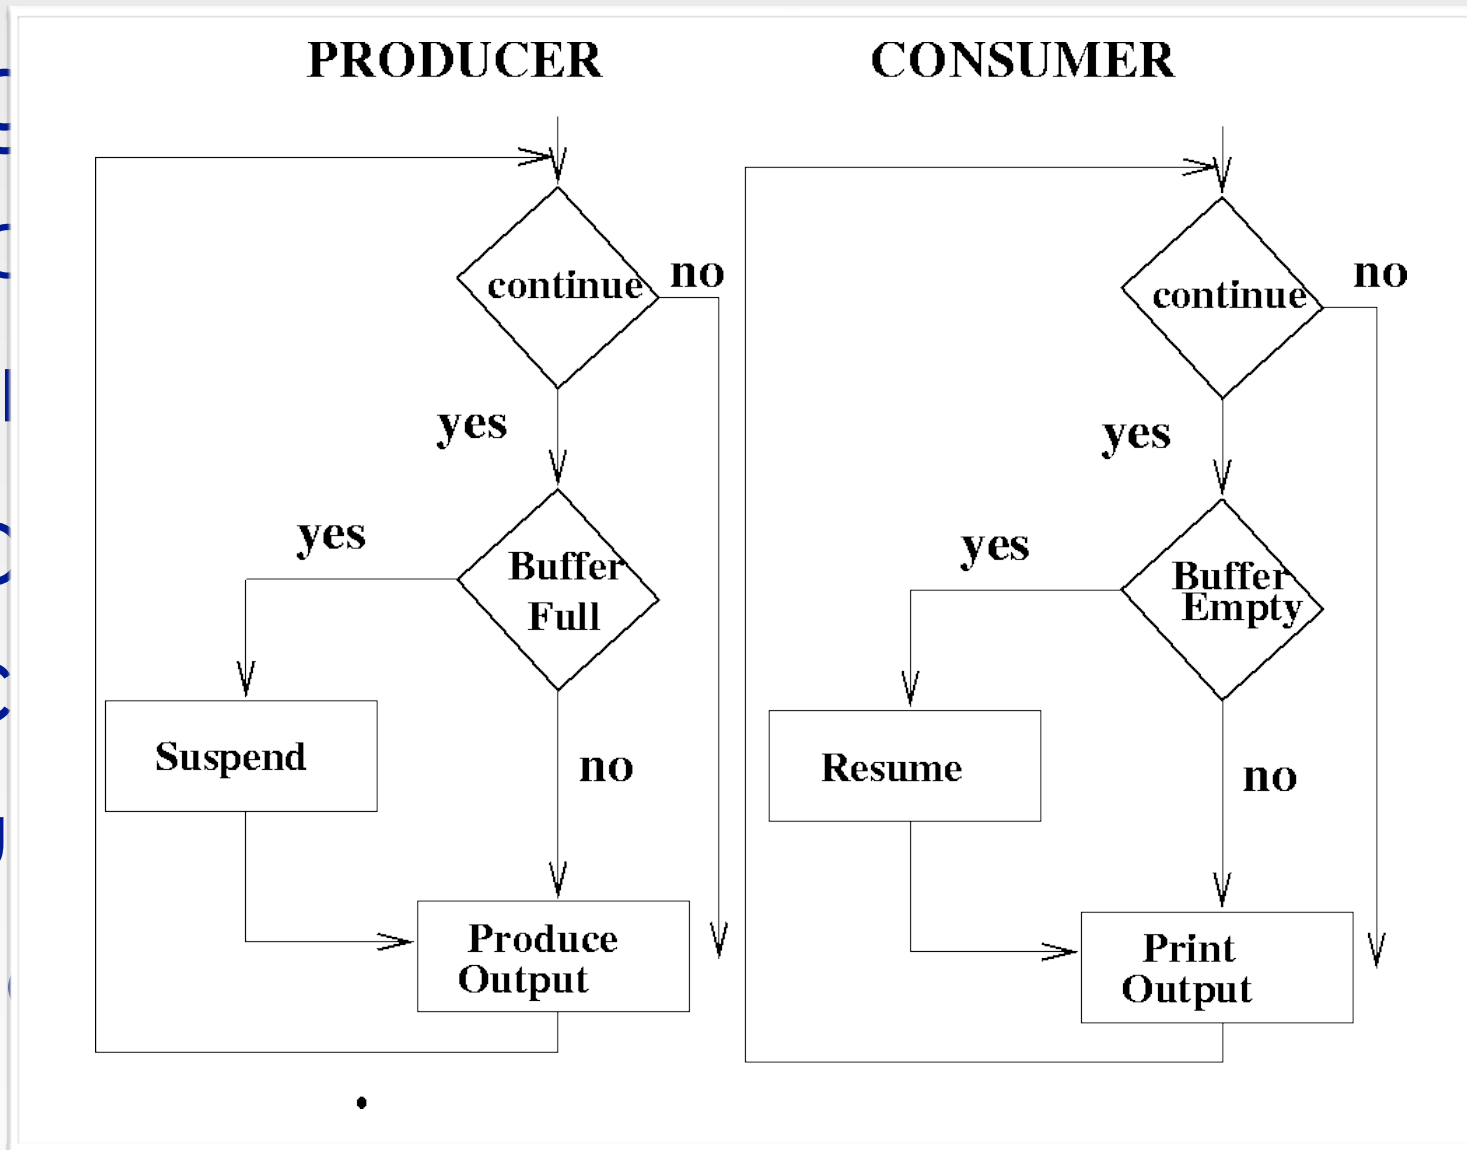
- Consumer

- Need to
produce

- Unbounded

- Bounded

- Producer and consumer must synchronize



Producer-Consumer Problem (cont.)

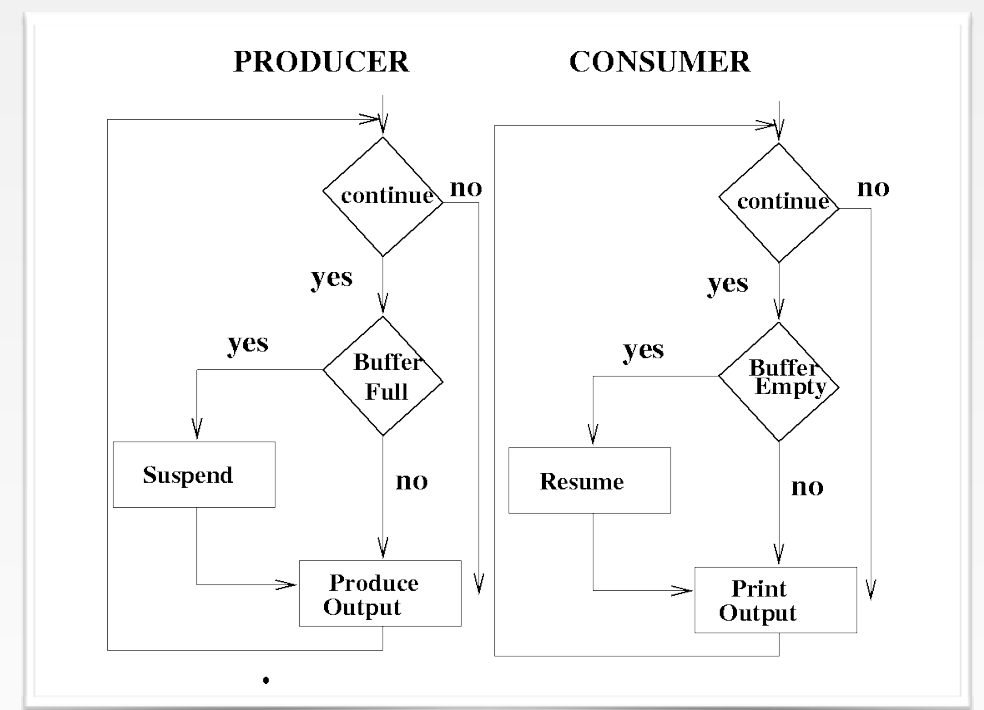
- Bounded-buffer using IPC (MP)

Producer

```
item next_produced;  
while (true){  
    /*produce an item in next  
    produced*/  
    send (next_produced)  
}
```

Consumer

```
item next_consumed;  
while (true){  
    while (in == out)  
        receive(next_consumed) ;  
    /*consume the item in the  
    next consumed*/  
}
```

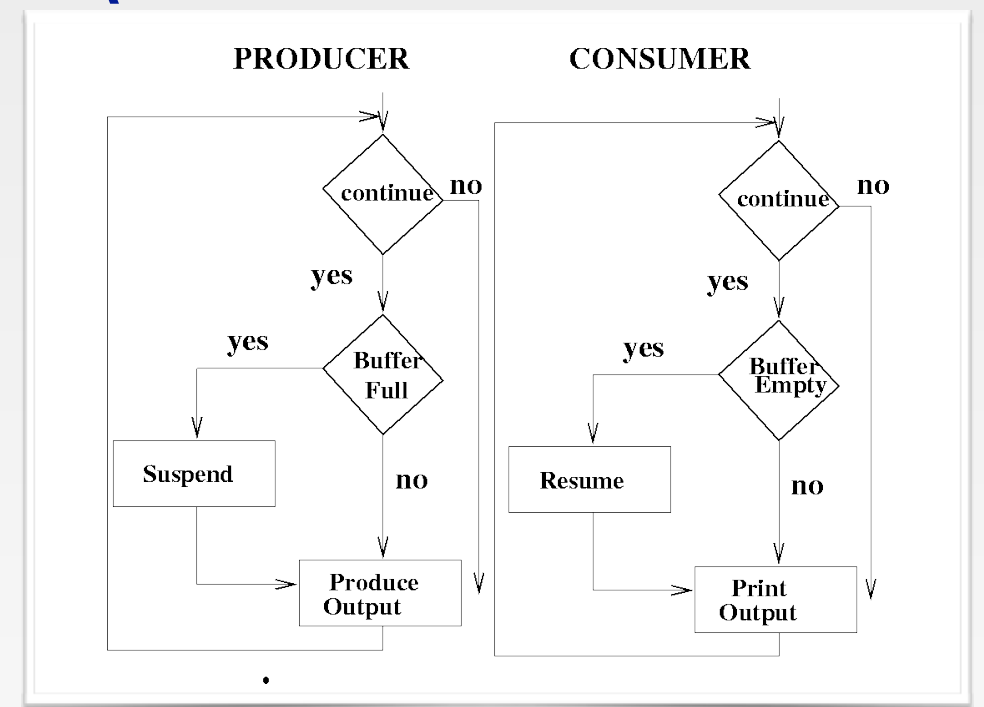


Producer-Consumer Problem (cont.)

- Bounded-buffer using IPC (shared memory solution)
- Shared data

```
#define BUFFER_SIZE 10
typedef struct{
...
}item;
```

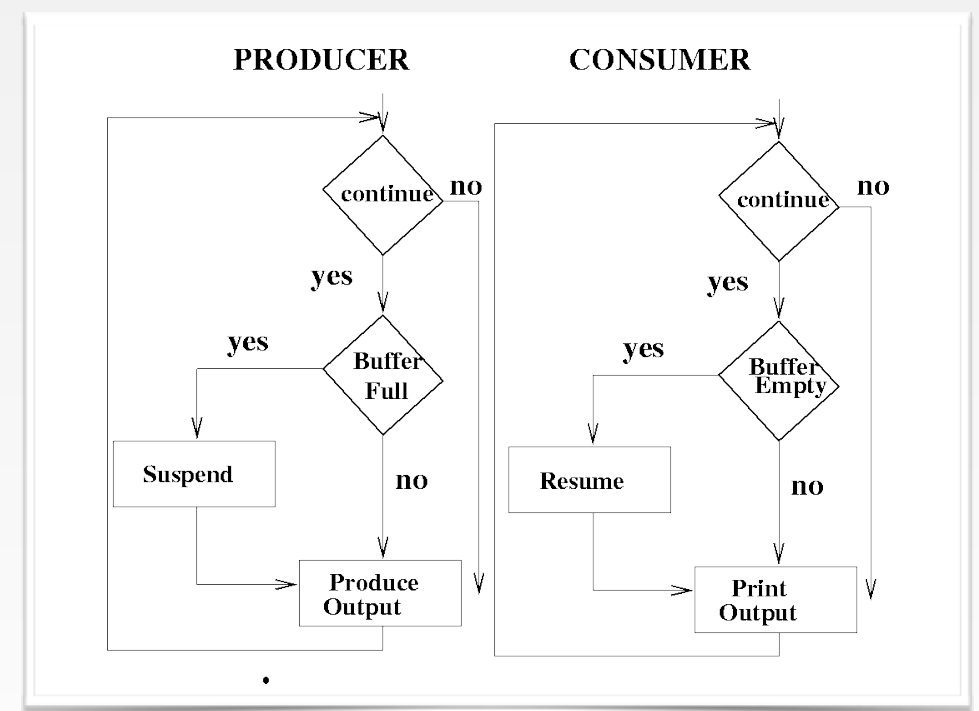
```
item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```



Producer-Consumer Problem (cont.)

- Bounded-buffer using IPC (shared memory solution)
- Producer- creates filled buffer

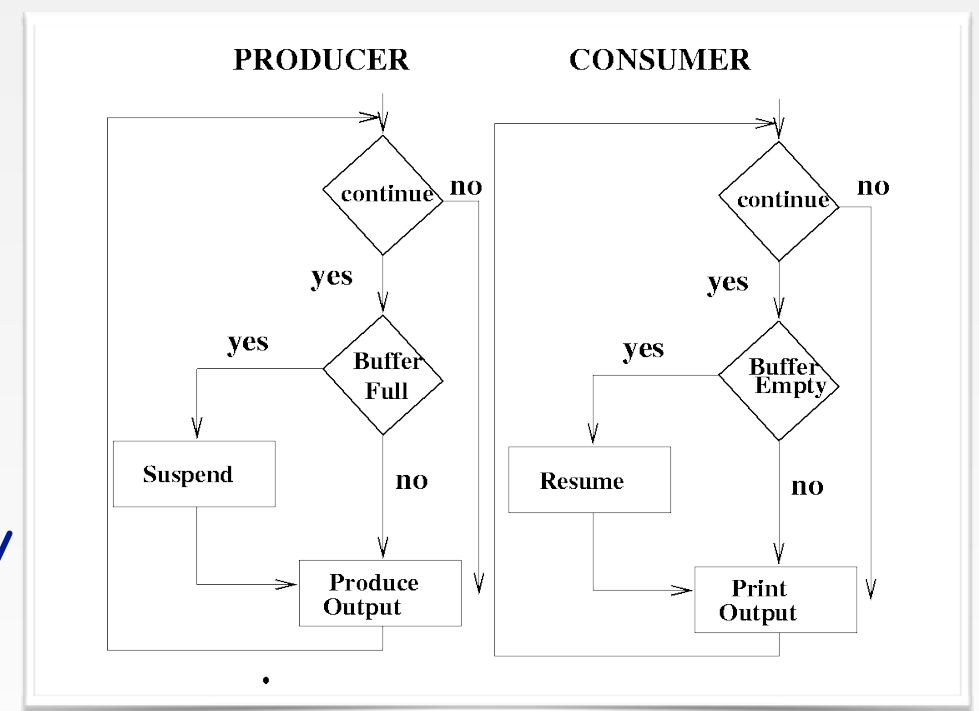
```
item next_produced;  
while (true){  
    /*produce an item in next produced*/  
    while (((in + 1) % BUFFER_SIZE) == out)  
        ; /*do nothing*/  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
}
```



Producer-Consumer Problem (cont.)

- Bounded-buffer using IPC (shared memory solution)
- Consumer- empties filled buffer

```
item next_consumed;  
while (true){  
    while (in == out)  
        /*do nothing*/  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    /*consume the item in the next consumed*/  
}
```



Shared Data

- Concurrent access to shared data may result in data inconsistency
- Data consistency requires orderly execution
- Shared memory solution allows at most $(n-1)$ items in the buffer at the same item



Bounded Buffer

- A solution that uses all N buffers is not that simple
 - Adding a variable *counter*
 - Initialized to 0
 - Incremented each time a new item is added



Bounded Buffer

- Shared data

```
#define BUFFER_SIZE 1-  
typedef struct{  
...  
}item;  
  
item buffer[BUFFER_SIZE];  
int in = 0;  
int out = 0;  
int counter = 0;
```



Bounded Buffer

- Producer process

```
while (true){  
    /*produce an item in the next  
    produced*/  
    while (counter == BUFFER_SIZE)  
        ; /*do nothing*/  
    buffer[in] = next_produced;  
    in = (in+1) % BUFFER_SIZE;  
    counter ++;  
    until false;  
}
```

- Consumer process

```
while (true){  
    while (counter == 0)  
        ; /*do nothing*/  
    next_consumed = buffer[out];  
    out = (out+1) % BUFFER_SIZE;  
    counter --;  
    ...  
    /* consume the item in next  
    consumed */  
}
```

The Statements must be executed **ATOMICALLY**



Problems is at the lowest level

- If threads are working on separate data, scheduling doesn't matter:

Thread A: $x = 1$; Thread B: $y = 2$;

- However, what about (initially, $y=12$)

Thread A: $x = 1$; $x = y+1$;

Thread B: $y = 2$; $y = y*2$;

- Or, what are the possible values of x

Thread A: $x = 1$; Thread B: $x = 2$;



The Critical-Section Problem

- N processes all competing to use shared data
 - Structure of process P_i
 - Each process has a code segment (or critical section)
 - Shared Data is accessed in the critical section

repeat

 entry section /*enter critical section*/

 critical section /*access shared variables*/

 exit section /*leave critical section*/

 remainder section /*do other work*/

until false;



Critical-Section Problem (cont.)

- Problem: Have to ensure that...
 - One process is executing in its critical section
 - No other process is allowed to execute in its critical section



Critical-Section Problem (cont.)

- Solutions - 3 requirements
 - Mutual Exclusion
 - Progress
 - Bounded Waiting
 - Assume that each process executes at a non-zero speed
 - No assumption concerning relative speed of n processes



Critical-Section Problem (cont.)

- Solutions - Initial Attempt

- Only 2 processes, P_0 and P_1
- General structure of process P_i

```
repeat
    entry section /*enter critical section*/
        critical section /*access shared variables*/
    exit section /*leave critical section*/
        remainder section /*do other work*/
until false;
```

- Processes may share some common variables to synchronize their actions



Critical-Section Handling in OS

- Preemptive
 - Allows preemption of process when running in kernel mode
- Non-preemptive
 - Runs until exits kernel mode, blocks, or voluntarily yields CPU
 - Essentially free of race conditions in kernel mode



Critical-Section Problem: Algorithm 1

- Shared variables

```
var turn = 0; /*indicating whose turn it is to enter  
               its critical section*/  
turn = i; /*implies that process Pi is allowed to enter  
           its critical section*/
```

- Process P₀

```
repeat  
    while turn ≠ 0 do noop;  
    critical section  
    turn = 1; /*leave critical section*/  
    remainder section /*do other work*/  
until false;
```

Satisfies mutual exclusion, but not progress



Critical-Section Problem: Algorithm 2

- Shared variables

```
boolean flag[2]; /*initially flag[0]=flag[1]=false*/  
flag[i] = true; /*implies that process Pi is ready to enter its critical  
section*/
```

- Process P_i:

Can block indefinitely, but progress requirement not met.

```
repeat  
    flag[i] = true;  
    while (flag[j]) do noop;  
    critical section  
    flag[i] = false;  
    remainder section  
until false
```



Critical-Section Problem: Algorithm 3

- Shared variables

```
boolean flag[2]; /*initially flag[0]=flag[1]=false*/  
flag[i] = true; /*implies that process Pi is ready to enter its critical  
section*/
```

- Process P_i:

Does not satisfy mutual exclusion requirement.

```
repeat  
    while (flag[j]) do noop;  
    flag[i] = true;  
    critical section  
    flag[i] = false;  
    remainder section  
until false
```



Critical-Section Problem: Algorithm 4

- Shared variables

```
int turn;  
boolean flag[2]; /*initially flag[0]=flag[1]=false*/  
flag[i] = true; /*implies that process Pi is ready to enter its critical  
section*/
```

- Process P_i

```
repeat  
    flag[i] = true;  
    turn = i;  
    while (flag[j] && turn == j) do noop;  
        critical section  
    flag[i] = false;  
    remainder section  
until false
```

Meets all three requirements, solves the critical section problem for 2 processes



Peterson's solution

- Provable that the 3 critical-section requirements are met:

- Mutual exclusion is preserved

- P_i enters critical section only if:

- either `flag[j] = false` or `turn == i`

- Progress requirement is satisfied

- Bounded-waiting requirements is met

```
repeat
    flag[i] = true;
    turn = j;
    while (flag[j] && turn == j) do noop;
        critical section
    flag[i] = false;
        remainder section
until false
```



Bakery Algorithm

- Critical section for n processes
 - Before entering its critical section, process receives a number. (Holder of the smallest number enters critical section)
 - If process P_i and P_j receive the same number
 - if $i \leq j$, then P_i is served first
 - else P_j is served first
 - The numbering scheme always generates number in increasing order of enumeration
 - i.e. 1,2,3,3,3,4,4,5,5



Bakery Algorithm (cont.)

- Notation

- Lexicographic order (ticket #, process id #)

- $(a, b) < (c, d)$ if $(a < c)$ or if $(a == c)$ and $(b < d)$

- $\max(a_0, \dots, a_{n-1}) = k,$

- such that $k \geq a_i$ (for $i = 0, \dots, n-1$, there is k ,)

- Shared Data

- boolean array: **choosing[n]** (initialized to false)

- int array: **turn[n]** (initialized to 0)



Bakery Algorithm (cont.)

```
int turn[n];
boolean choosing[n];
int j;
while(1){
    choosing[i] = true;
    turn[i] = 1 + max(turn[0], turn[1], ... turn[n-1]);
    choosing[i] = false;
    for (j = 0; j < n; j++){
        if (j != i){
            //Wait until thread j receives its number:
            while (choosing[j]);
            //Wait until all threads with smaller numbers or with the same number
            //but with higher priority, finish their work:
            while (turn[j] != 0 && ((turn[j], j) < (turn[i], i)));
        }
        critical section;
        turn[i] = 0;
        non-critical section;
    }
}
```

/* notation: $(a, b) < (c, d)$ is equivalent to
 $(a < c) \ || \ (a = c \ \&\& \ b < d)$ */



Supporting Synchronization

Programs	Shared Programs
High-Level API	Locks, Semaphores, Monitors, Send/Receive, CCregions
Hardware	Load/Store, Disable Ints, Test/Set, Comp/Swap



HW Solutions for Sync.

- Load/Store

- Atomic operations required for synchronization
- Shows how to protect a critical section with only atomic load and store

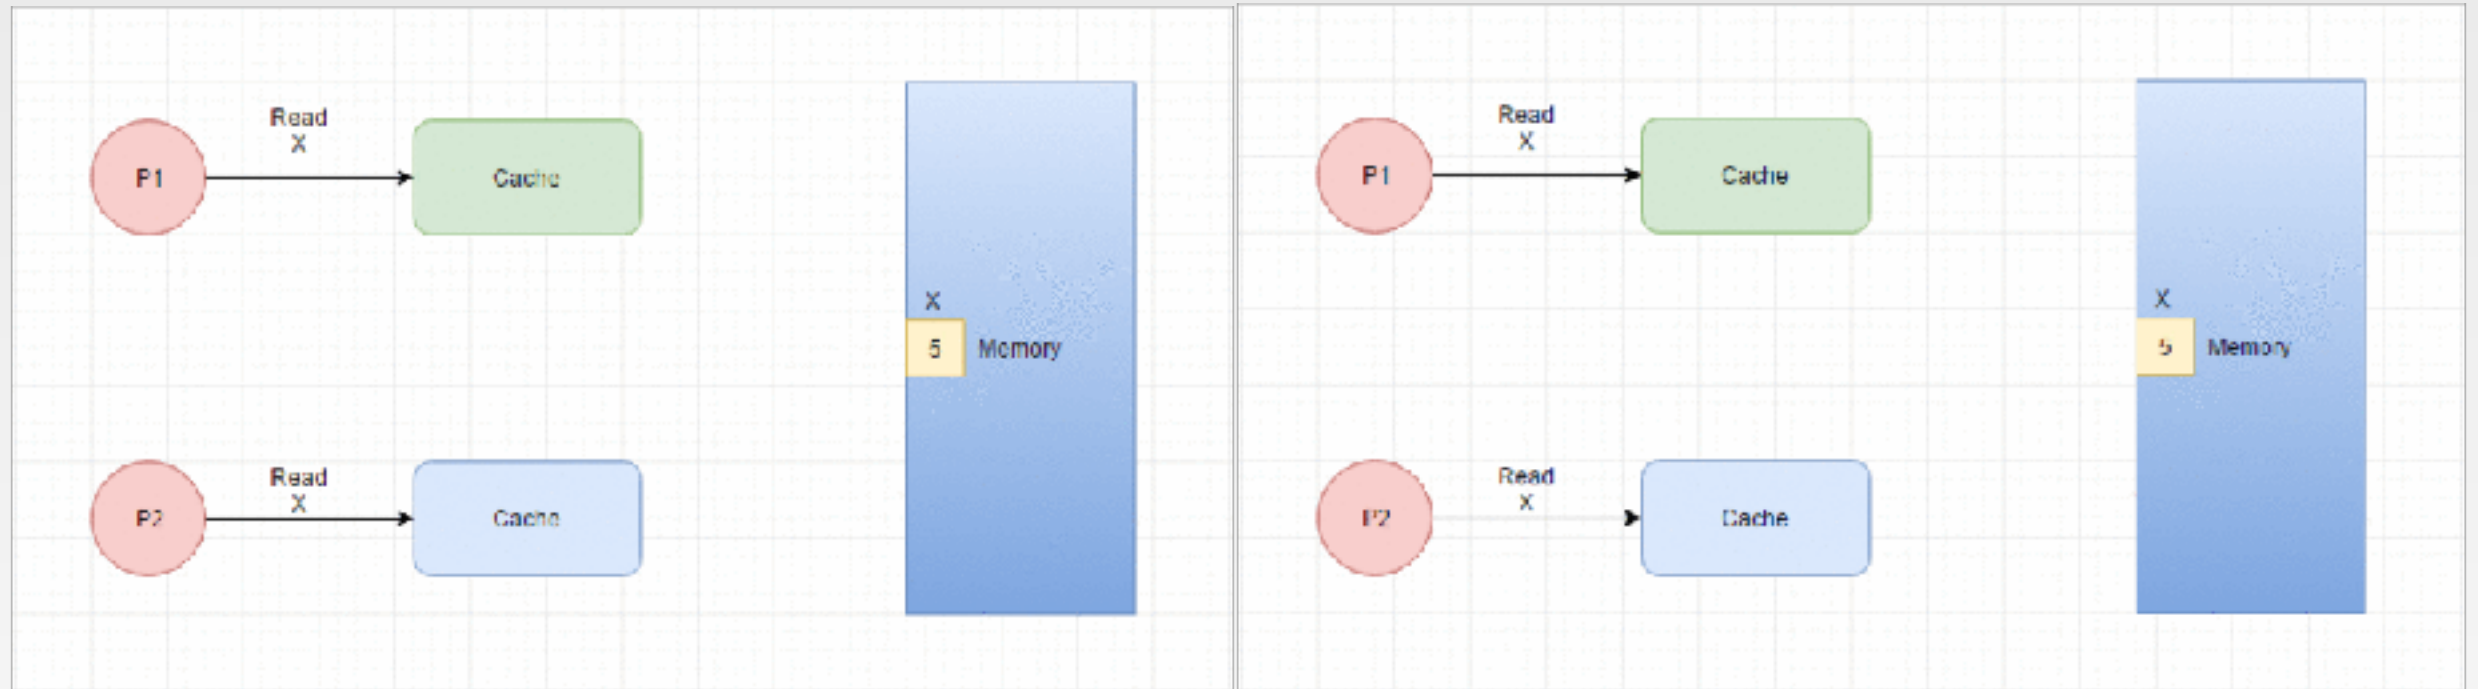


HW Solutions for Sync. (cont.)

- Mutual exclusion solutions presented depend on memory HW having R/W cycle
 - multiple R/W + same location *Would Not Work* the time
 - Processors with caches but no cache coherency can NOT use the solutions



Recall: cache coherence



No Cache Coherence

Cache Coherence

Synchronization Hardware

- Based on idea of locking
 - Protecting critical regions via locks
- Uniprocessors - could disable interrupts
 - Currently running code → execute w/o preemption
 - Too inefficient on multiprocessor systems
- Modern machine → atomic instructions
 - Either test memory word and set value
 - Or swap contents of two memory words



Locks

- Solutions to critical-section problem using Locks

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
}
```



test_and_set Instruction

- Test and modify the content of a word atomically - (**test_and_set** instruction)
- Similarly “SWAP” instruction

```
boolean test_and_set (boolean *target) {  
    boolean rv = *target;  
    *target = TRUE;  
    return rv;  
}
```
- Executed atomically
- Returns the original value of passed parameter
- Set the new value of passed parameter of “TRUE”



Mutual Exclusion with `test_and_set`

- Shared boolean “lock”, initialized to False
- Solution:

```
do {  
    while (test_and_set(&lock))  
        ; /*do nothing*/  
        /*critical section*/  
    lock = false;  
        /*remainder section*/  
}while (true)
```



compare_and_swap Instruction

- Definition

```
int compare_and_swap (int *value, int expected, int new_value) {  
    int temp = *value;  
    if (*value == expected)  
        *value = new_value;  
    return temp;  
}
```

- Executed atomically
- Returns the original value of passed “value”
- Swap takes place only if “value == expected”



Mutual Exclusion with `compare_and_swap`

- Shared int “lock” initialized to 0
- Solution:

```
do {  
    while (compare_and_swap(&lock, 0, 1) != 0)  
        ; /*do nothing*/  
        /*critical section*/  
    lock = 0;  
        /*remainder section*/  
}while (true)
```



Bounded Mutual Exclusion with test_and_set

```
do {
    boolean waiting[i] = true;
    boolean key = true;
    while (waiting[i] && key)
        key = test_and_set (&lock);
    waiting[i] = false;
    /*critical section*/
    j = (i+1) % n;
    while ((j != i) && !waiting[j])
        j = (j + 1) % n;
    if (j == i)
        lock = false;
    else
        waiting[j] = false;
    /*remainder section*/
}while (true)
```



Mutex Locks

- SW tool to solve Critical-Section problem
- Simplest is Mutex Lock
 - First **acquire()** a lock
 - Then **release()** the lock
 - Boolean indicating if lock is on or not
- Calls must be Atomic
 - Usually via HW atomic instructions
- This solution requires busy waiting
 - spinlock



Mutex Locks (cont.)

```
    acquire {  
        while (!available)  
            ; /*busy wait*/  
        available = false;  
    }  
  
    release {  
        available = true;  
    }  
  
do {  
    acquire lock;  
    /*critical section*/  
    release lock;  
    /*remainder section*/  
}while (true)
```



HW Support: Other Examples

```
swap (&address, register){ /*86*/
    temp = M[address];
    M[address] = register;
    register = temp;
}while (true)
```


```
load-linked&store conditional
(&address){ /*R4000, alpha*/
    loop:
        ll r1, M[address];
        /*Can do arbitrary comp */
        move r2, 1;
        sc r2, M[address];
        begz r2, loop;
    }
}
```

```
compare&swap (&address, reg1, reg2){
/*68000*/
    if (reg1 == M[address]){
        M[address] = reg2;
        return success;
    } else{
        return failure;
    }
}
```



Semaphore

- Synchronization tool
- More sophisticated way than Mutex Lock
- Semaphore S : int variable
- Can only be accessed via two atomic operations
 - **wait()** and **signal()**
 - Originally called **P()** and **V()**



Semaphore (cont.)

- Definition of the **wait()** operation

```
wait(S) {  
    while (S <= 0)  
        ; /*busy wait*/  
    S--;  
}
```

- Definition of the **signal()** operation

```
signal(S) {  
    S++;  
}
```



Semaphore (cont.)

- Usage
 - Counting semaphore
 - int range over an unrestricted domain
 - Binary semaphore
 - int range only between 0 and 1
 - Same as a Mutex Lock



Semaphore (cont.)

- Usage (cont.)

- Consider P_1 and P_2 that require S_1 to happen before S_2
- Create a semaphore “synch” initialized to 0

```
P1:  
    S1;  
    signal (synch) ;  
P2:  
    wait (synch) ;  
    S2;  
}
```

- Can implement a counting semaphore S as a binary semaphore



Semaphore: Problem...

- Locks prevent conflicting actions on shared data
 - Lock before entering critical section
 - Lock before accessing shared data
 - Unlock when leaving, after accessing shared data
 - Wait if locked



Semaphore: Problem... (cont.)

- All synchronization involves waiting
 - Busy Waiting (spinlock)
 - Waiting thread may take cycles away from thread holding lock
 - OK for short time (prevents context switch)
 - Priority Inversion
- For longer runtimes, need to modify P and V so that processes can *block* and *resume*



Semaphore: Implementation

- Must guarantee...NOT on the same semaphore at the same time
- Thus, **wait** and **signal** code are in placed in the critical section
 - Can have busy waiting in C-S implementation
 - But implementation code is short
 - Little busy waiting if C-S rarely occupied
- Apps may spend lot of time in C-S,
so no a good solution



Semaphore: Implementation (cont.)

- Two operations
 - *block* - suspends the process that invokes it
 - *wakeup* - resumes the execution of a blocked process P



Semaphore: Implementation (cont.)

- With each semaphore there is an associated waiting queue
- Each entry in a waiting queue has two data items:
 - Value (type: integer)
 - Pointer to next record in the list

```
typedef struct{  
    int value;  
    struct process *list;  
} semaphore
```



Semaphore: Implementation (cont.)

- Semaphore operations are now defined as:

```
wait (semaphore *S) {
    S->value --;
    if (S->value < 0) {
        add this process to S->list;
        block();
    }
}

signal (semaphore *S) {
    S->value ++;
    if (S->value <= 0) {
        remove a process P from S->list;
        wakeup(P);
    }
}
```



Block/Resume Semaphore

- If process is blocked, enqueue PCB of process and call scheduler to run a different process



Block/Resume Semaphore (cont.)

- Semaphores are executed atomically
 - No two processes execute *wait* and *signal* at the same time
 - Mutex can be used to make sure that two processes do not change count at the same time
 - If an interrupt occurs while mutex is held, it will result in a long delay
 - Solution: Turn off interrupts during critical section



Deadlock and Starvation

- Deadlock

- Two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes

P_0	P_1
<code>wait (S) ;</code>	<code>wait (Q) ;</code>
<code>wait (Q) ;</code>	<code>wait (S) ;</code>
...	...
<code>signal (S) ;</code>	<code>signal (Q) ;</code>
<code>signal (Q) ;</code>	<code>signal (S) ;</code>



Deadlock and Starvation (cont.)

- Starvation
 - Indefinite blocking
 - A process may never be removed from the semaphore queue in which it is suspended
- Priority Inversion
 - Scheduling problem
 - Lower-priority process holds a lock needed by higher-priority process
 - Solved via priority-inheritance protocol



Classical Problems of Sync.

- Bounded-Buffer problem
- Readers and Writers problem
- Dining-Philosophers problem



Bounded-Buffer Problem

- n buffers, one item each
- Semaphore mutex (initialized to 1)
- Semaphore full (initialized to 0)
- Semaphore empty (initialized to n)



Bounded-Buffer Problem (cont.)

- Producer process (creates filled buffers)

```
do {  
    ...  
    /*produce an item in next_produced*/  
    ...  
    wait (empty);  
    wait (mutex);  
    ...  
    /*add next produced to the buffer*/  
    ...  
    signal(mutex);  
    signal(full);  
} while (true);
```



Bounded-Buffer Problem (cont.)

- Consumer process (empties filled buffers)

```
do {  
    wait (full);  
    wait (mutex);  
    ...  
    /*remove an item from to the next_consumed*/  
    ...  
    signal(mutex);  
    signal(empty);  
    ...  
    /*consume the item in next consumed*/  
    ...  
} while (true);
```



Discussion

- Asymmetry?
 - Producer does: **wait(empty)** , **signal(full)**
 - Consumer does: **wait(full)** , **signal(empty)**
- Is order of wait()'s important?
 - Yes! Can cause deadlock
- Is order of signal()'s important?
 - No, except that it might affect scheduling efficiency



Readers/Writers Problem


- Motivation: shared database

- Two classes of users
 - Readers - no modify
 - Writers - read&modify

- Is using a single lock

on the whole database sufficient?

- Like to have many readers at the same time
- Only one writer at a time



Readers-Writers Problem (cont.)

- Shared Data
 - Data set
 - Semaphore **rw_mutex** initialized to 1
 - Semaphore **mutex** initialized to 1
 - Int **read_count** initialized to 0



Readers-Writers Problem (cont.)

- The structure of a writer process

```
do {  
    wait (rw_mutex);  
    ...  
    /*writing is performed*/  
    ...  
    signal(rw_mutex);  
} while (true);
```



Readers-Writers Problem (cont.)

- The structure of a reader process

```
do {  
    wait (mutex);  
    read_count++;  
    if (read_count == 1)  
        wait (rw_mutex);  
    signal(mutex);  
    ...  
    /*reading is performed*/  
    ...  
    wait (mutex);  
    read count --;  
    if (read_count == 0)  
        signal (rw_mutex);  
    signal (mutex);  
} while (true);
```



Readers-Writers Problem Variations

1. No reader kept waiting unless writer has permission to use shared object
2. Once writer is ready, it performs the write ASAP



Dining-Philosophers Problem

- Either thinking or eating
- Don't interact with neighbors
- Pick up 2 chopsticks to eat
 - One at a time
 - Need both to eat
 - Then release both when done
- Shared Data
 - Bowl of rice (data set)
 - Semaphore **chopstick[5]** (init to 1)



Dining-Philosopher Problem (cont.)

- The structure of Philosopher i:

```
do {  
    wait (chopstick[i]);  
    wait (chopstick[(i+1) % 5]);  
    //eat  
    signal (chopstick[i]);  
    signal (chopstick [(i+1) % 5]);  
    //think  
    ...  
} while (true);
```



- What's the problem with this algorithm?

Dining-Philosopher Problem (cont.)

- What's the problem?
 - Deadlock handling
 - Allow at most 4 philosophers to be sitting simultaneously at the table
 - Allow a philosopher to pick up the chopsticks only if both are available (picking must be done in a critical section)
 - Use an asymmetric solution



Higher Level Synchronization

- Timing errors are still possible with semaphores
 - `signal(mutex) ... wait(mutex)`
 - `wait(mutex) ... wait(mutex)`
 - `wait(mutex) ... /*forget to signal*/`
- Deadlock and starvation are possible



Motivation for Other Sync. Constructs

- Semaphores are a huge step up from loads and stores
 - Problem is that semaphores are dual purpose
 - Used for both mutex and scheduling constraints
 - E.g.: The fact that flipping of wait()'s in bounded buffer gives deadlock is not immediately obvious. How do you prove the correctness to someone?



Motivation for Other Sync. (cont.)

- Idea: allow manipulation of a shared variable only when condition is met
 - Conditional critical region
- Idea: use locks for mutual exclusion and condition variable for scheduling constraints
 - Monitor



Conditional Critical Regions

- High-level synchronization construct
- Shared variable v of type T
 - **var v : shared T**
- Variable v is accessed only inside statement
 - **region v when B do S**
 - B : boolean expression
 - No other process can access v while S is being executed



Critical Regions (cont.)

- Region referring to the same shared variable exclude each other in time
- When a process tries to execute the region statement
 - If B is true, S is executed
 - If B is false, the process is delayed until B becomes true, and no other process is in the region associated with v




```


1  region v when B do S() {
2      semaphore xMutex, // mutual exclusion
3          xDelay; //
4      int xCount, // number of processes waiting for access to region v
5          xTemp; // number of processes allowed to check condition B
6
7      p(xMutex);
8      if(!B) { // not B? we have to wait until it is B
9          xCount++; // we're also waiting for not B to become B
10         v(xMutex); // wait until you can go further
11         p(xDelay); //
12         while(!B) { //
13             xTemp++; //
14             if(xTemp < xCount) v(xDelay); //
15             else v(xMutex); //
16             v(xDelay); //
17         }
18         xCount--; // got out of while(!B), it means B is true so we're not waiting anymore
19     }
20     S(); // execute sequence of instructions
21     if(xCount > 0) { // if someone's waiting for B...
22         xTemp = 0; //
23         v(xDelay); //
24     } else v(xMutex); // if noone wants B, simply give up the right to the critical region
25 }

```

Monitors

- High-level synchronization construct that allows the safe sharing of an abstract data type among concurrent processes

```
monitor monitor_name {  
    /* shared variable declare*/  
    procedure P1 (...) {...};  
  
    procedure Pn (...) {...};  
  
    initialization code (...) {...}  
    ...  
}
```



Monitor with Condition Variables

- Lock: provides mutual exclusion to share data
 - Always require before accessing shared data structure
 - Always release after finishing with shared data
 - Lock initially free
- Condition Variable: threads waiting for something inside a critical section
 - Key idea: go to sleep, automatically releasing lock at time going to sleep




Condition Variables

- **condition x, y ;**
- Two operations are allowed on a condition variable:
 - **$x.wait()$**
 - that invoking this operation is suspended until another process invokes **$x.signal()$**
 - **$x.signal()$**
 - resume exactly one suspended process
 - If no process is suspended, then the signal operation has no effect



Monitor with Condition Variables



Monitor with Condition Variables

- Condition variables Choices
 - If P invokes **`x.signal()`**, and Q is suspended in **`x.wait()`**, what should happen next?
 - Both Q and P cannot execute in parallel.
 - If Q is resumed, the P must wait
 - Operation includes
 - Signal and wait - P waits until Q either leaves the monitor or it waits for another condition
 - Signal and continue - Q waits until P leaves the monitor or it waits for another condition



```

monitor DiningPhilosophers {
    enum {THINKING, HUNGRY, EATING} state[5];
    condition self[5];
    void pickup (int i){
        state[i] = HUNGRY;
        test(i); /*test left and right are not eating*/
        if (state[i] != EATING) self[i].wait;
    }
    void putdown (int i){
        state[i] = THINKING;
        /*test left and right neighbors*/
        test ((i + 4) % 5); /*signal on neighbor*/
        test ((i + 1) % 5); /*signal other neighbor*/
    }
    void test (int i) {
        if ((state[(i + 4) % 5] != EATING)) &&
            (state[i] == HUNGRY)&&
            (state[(i + 1) % 5] != EATING) {
                state[i] = EATING;
                self.signal();
            }
    }
    initialization_code() {
        for (int i = 0; i < 5; i++)
            state[i] = THINKING;
    }
}

```

Monitor Solution to Dining Philosophers



Monitor Solution to Dining Philosophers (cont.)

- Each philosopher i invokes the operation **pickup()** and **putdown()** in the following sequence

`DiningPhilosophers.pickup(i) ;`

EAT

`DiningPhilosophers.putdown(i) ;`

- No deadlock, but starvation is possible



Monitor Implementation using Semaphore

- Variables

```
semaphore mutex; /*initially = 1*/  
semaphore next;  /*initially = 0*/  
int next_count = 0;
```
- Each procedure F will be replaced by

```
wait(mutex) ;  
  
...  
body of F;  
  
...  
if (next_count) > 0  
    signal (next);  
else  
    signal (mutex);
```
- Mutual Exclusion is ensured



Monitor Implementation using Conditional Variables

- For each condition variable x , we have

```
semaphore x_sem; /*initially = 0*/  
int x_count = 0;
```

- The operation **$x.\text{wait}$** can be implemented as

```
x_count ++;  
if (next_count > 0) {  
    signal (next);}  
else {  
    signal (mutex);}  
wait (x_sem);  
x_count --;
```



Monitor Implementation using Conditional Variables (cont.)

- The operation **`x.signal`** can be implemented as

```
if (x_count > 0) {  
    next_count ++;  
    signal (x_sem);  
    wait (next);  
    next_count --;  
}
```



Resuming Processes within Monitor

- If several processes queued on condition ***x***, and ***x.signal()*** executed, which should be resumed?
- FCFS frequently not adequate
- *conditional-wait* construct of the form ***x.wait(c)***
 - ***c*** is priority number
 - Processes with lowest number (highest priority) is scheduled next



Single Resource Allocation

- a single resource
- among competing processes
- using priority numbers

```
R.acquire (t);  
...  
access the resource;  
...  
R.release;
```

- R: an instance of type **ResourceAllocator**



Monitor to Allocate Single Resource

```
monitor ResourceAllocator {
    boolean busy;
    condition x;
    void acquire (int time) {
        if (busy)
            x.wait(time);
        busy = TRUE;
    }
    void release () {
        busy = FALSE;
        x.signal ();
    }
    initialization code() {
        busy = FALSE;
    }
}
```



Summary

- Why synchronization
 - cooperating sequential processes
 - share data
 - must provide mutual exclusion
 - critical section code used by only one process/thread at a time
- Synchronization problems
 - bounded-buffer/ reader-writer/ dining-philosopher
- Synchronization implementation
 - Monitor

